



中国科学院大学
University of Chinese Academy of Sciences

Introduction to OpenMP

Part3. *Working with OpenMP*

人工智能学院
缪青海

miaoqh@ucas.ac.cn



Content

- Getting started with OpenMP
- Core features of OpenMP
- [Work with OpenMP](#)
- Advanced topics on OpenMP



Working with OpenMP

- Synchronization 线程同步
- Data environment 数据属性
- Linked lists traversal 链表遍历



Synchronization: Barrier

- **Barrier:** Each thread waits until all threads arrive.

```
#pragma omp parallel shared (A, B, C) private(id)
{
    id=omp_get_thread_num();
    A[id] = big_calc1(id);
    #pragma omp barrier
    #pragma omp for
        for(i=0;i<N;i++){C[i]=big_calc3(i,A);}
    #pragma omp for nowait
        for(i=0;i<N;i++){ B[i]=big_calc2(C, i); }
    A[id] = big_calc4(id);
}
```

implicit barrier at the end of a
for worksharing construct

implicit barrier at the end
of a parallel region

no implicit barrier
due to nowait



Master Construct

- The **master** construct denotes a structured block that is only executed by the **master thread**.
- The other threads just skip it (no synchronization is implied).

master默认
没有barrier

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp master
    { exchange_boundaries(); }
    #pragma omp barrier
    do_many_other_things();
}
```



Single Construct

- The **single** construct denotes a block of code that is executed by only one thread (**not necessarily the master thread**).
- A barrier is implied at the end of the single block (can remove the barrier with a **nowait** clause).

与master区别:
默认有barrier

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp single
    { exchange_boundaries(); }
    do_many_other_things();
}
```



Sections Construct

- The **Sections** worksharing construct gives a different structured block to each thread.
- By default, there is a barrier at the end of the “omp sections”. Use the “**nowait**” clause to turn off the barrier.

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        x_calculation();
        #pragma omp section
        y_calculation();
        #pragma omp section
        z_calculation();
    }
}
```



Synchronization: Lock routines

- A lock implies a memory fence (a “flush”) of all thread visible variables.
- Simple Lock: 比critical等更底层
 - A simple lock is available if it is unset.
 - `-omp_init_lock()`, `omp_set_lock()`, `omp_unset_lock()`,
`omp_test_lock()`, `omp_destroy_lock()`
- Nested Locks:



Synchronization: Lock routines

- A lock implies a memory fence (a “flush”) of all thread visible variables.
- Simple Lock:
- Nested Locks:
 - A nested lock is available if it is unset or if it is set but owned by the thread executing the nested lock function
 - `-omp_init_nest_lock()`, `omp_set_nest_lock()`,
`omp_unset_nest_lock()`, `omp_test_nest_lock()`,
`omp_destroy_nest_lock()`




Synchronization: Simple Lock

```
#pragma omp parallel for  
for(i=0;i<NBUCKETS; i++){  
    omp_init_lock(&hist_locks[i]);    hist[i] = 0;  
}
```

One lock per element of hist

```
#pragma omp parallel for  
for(i=0;i<NVALS;i++){  
    ival = (int) sample(arr[i]);  
    omp_set_lock(&hist_locks[ival]);  
    hist[ival]++;  
    omp_unset_lock(&hist_locks[ival]);  
}
```



Enforce mutual exclusion on update to hist array

```
for(i=0;i<NBUCKETS; i++)  
    omp_destroy_lock(&hist_locks[i]);
```

Free-up storage when done.



Runtime Library routines

■ Runtime environment routines – 1 :

☐ Modify/Check the number of threads

- `omp_set_num_threads()`, `omp_get_num_threads()`,
`omp_get_thread_num()`, `omp_get_max_threads()`

☐ Are we in an active parallel region?

- `omp_in_parallel()`



Runtime Library routines

■ Runtime environment routines – 2 :

- ☐ Do you want the system to dynamically vary the number of threads from one parallel construct to another?
 - `omp_set_dynamic()`, `omp_get_dynamic()`
- ☐ How many processors in the system?
 - `omp_num_procs()`



Runtime Library routines

- Use fixed number of threads in a program:

```
#include <omp.h>
void main()
{   int num_threads;
    omp_set_dynamic( 0 );
    omp_set_num_threads( omp_num_procs() );
    #pragma omp parallel
    {   int id=omp_get_thread_num();
        #pragma omp single
        num_threads = omp_get_num_threads();
        do_lots_of_stuff(id);
    }
}
```



Environment Variables

- Set the default number of threads to use.
 - `-OMP_NUM_THREADS int_literal`
- Control the size of child threads' stack
 - `-OMP_STACKSIZE`
- Hint to runtime how to treat idle threads
 - `-OMP_WAIT_POLICY`
 - `-ACTIVE` keep threads alive at barriers/locks
 - `-PASSIVE` try to release processor at barriers/locks
- Bind threads to processor
 - `-OMP_PROC_BIND`
 - `-TRUE / FALSE` turn on for better use of caches



Working with OpenMP

- Synchronization 线程同步
- Data environment 数据属性
- Linked lists traversal 链表遍历



Data environment

■ Default storage attributes:

□ Shared Memory programming model:

- Most variables are shared by default (**heap**)

□ Global variables are SHARED among threads:

- File scope variables, static
- dynamically allocated memory (ALLOCATE, malloc, new)

■ But not everything is shared...

- **Stack** variables in functions called from parallel regions are PRIVATE
- Automatic variables within a statement block are PRIVATE.

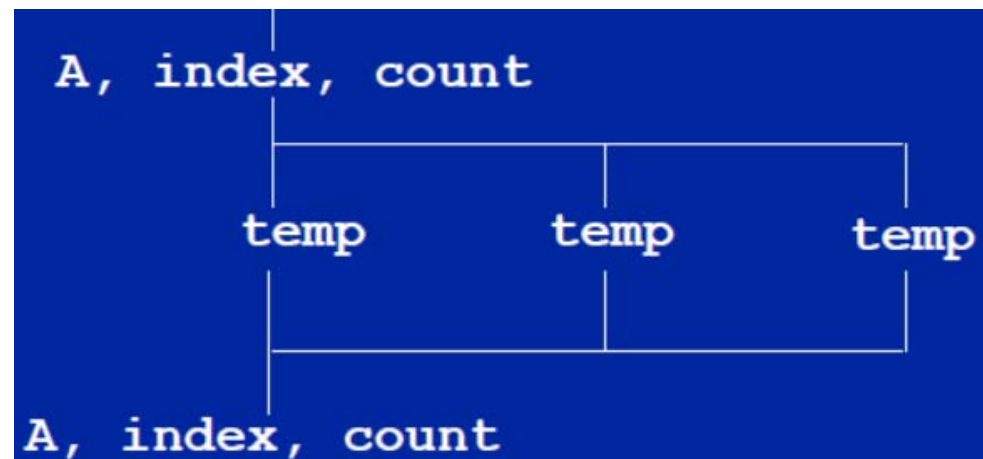


Data environment: example

```
double A[10];
int main() {
    int index[10];
    #pragma omp parallel
        work(index);
    printf("%d\n", index[0]);
}
```

```
extern double A[10];
void work(int *index) {
    double temp[10];
    static int count;
    ...
}
```

- **A, index and count** are shared by all threads. (**heap**)
- **temp** is local to each thread. (**stack**)





Changing storage attributes

- **One can selectively change storage attributes** for constructs using the following clauses:
 - – SHARED
 - – PRIVATE
 - – FIRSTPRIVATE
- The final value of a private inside a parallel loop can be transmitted to the shared variable outside the loop with:
 - – LASTPRIVATE
- The default attributes can be overridden with:
 - – DEFAULT (SHARED | NONE)



Data Sharing: Private Clause

■ `private(var)`

- creates a new local copy of `var` for each thread.
- – The value of the private copies is **uninitialized**
- – The value of the original variable is **unchanged after the region**

```
void wrong() {  
    int tmp = 0;  
    #pragma omp parallel for private(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j;  
    printf("%d\n", tmp);  
}
```



Data Sharing: Firstprivate Clause

■ **firstprivate(var)**

- Variables initialized from shared variable
- C++ objects are copy-constructed

```
incr = 0;  
#pragma omp parallel for firstprivate(incr)  
for (i = 0; i <= MAX; i++) {  
    if ((i%2)==0) incr++;  
    A[i] = incr;  
}
```

Each thread gets its own copy
of incr with an initial value of 0



Data Sharing: Lastprivate Clause

■ `lastprivate(var)`

- Variables update shared variable using value from last iteration
- C++ objects are updated as if by assignment

```
void sq2(int n, double *lastterm) {  
    double x; int i;  
    #pragma omp parallel for lastprivate(x)  
        for (i = 0; i < n; i++){  
            x = a[i]*a[i] + b[i]*b[i];  
            b[i] = sqrt(x);  
        }  
    *lastterm = x;  
}
```

“x” has the value it held
for the “last sequential”
iteration (i.e., for $i=(n-1)$)



Data Sharing: A Test

```
variables: A = 1, B = 1, C = 1  
#pragma omp parallel private(B) firstprivate(C)
```

- Are A,B,C local to each thread or shared inside the parallel region?
- What are their initial values inside and values after the parallel region?



Data Sharing: A Test

```
variables: A = 1, B = 1, C = 1  
#pragma omp parallel private(B) firstprivate(C)
```

Inside this parallel region ...

- “A” is shared by all threads; equals 1
- “B” and “C” are local to each thread.
 - B’s initial value is undefined
 - C’s initial value equals 1

Following the parallel region ...

- B and C revert to their original values of 1
- A is either 1 or the value it was set to inside the parallel region



Data Sharing: Default Clause

■ DEFAULT(SHARED)

- ☐ Note this is the default storage attribute
- ☐ Exception: `#pragma omp task`

■ DEFAULT(PRIVATE)

- ☐ To change default to PRIVATE
- ☐ *each* variable in the construct is made private as if specified in a private clause

■ DEFAULT(NONE):

- ☐ no default for variables in static extent. Must list storage attribute for each variable in static extent.
- ☐ Good programming practice!



Example: Serial PI Program

```
static long num_steps = 100000;
double step;
int main ()
{
    int i; double x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;

    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```



Example: PI Program minimal changes

```
#include <omp.h>
static long num_steps = 100000; double step;
int main ()
{
    int i; double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel for private(x) reduction(+:sum)
    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```



Working with OpenMP

- Synchronization 线程同步
- Data environment 数据属性
- Linked lists traversal 链表遍历



Review

- To create a team of threads
 - ☐ ☐ `#pragma omp parallel`
- To share work between threads:
 - ☐ ☐ `#pragma omp for`
 - ☐ ☐ `#pragma omp single`
- To prevent conflicts (prevent races)
 - ☐ ☐ `#pragma omp critical`
 - ☐ ☐ `#pragma omp atomic`
 - ☐ ☐ `#pragma omp barrier`
 - ☐ ☐ `#pragma omp master`



Review

■ Data environment clauses

- □ `private (variable_list)`
- □ `firstprivate (variable_list)`
- □ `lastprivate (variable_list)`
- □ `reduction(+:variable_list)`

- Where `variable_list` is a comma separated list of variables.



Simple list traversal

- Given what we've covered about OpenMP, how would you process this loop in Parallel?

```
p=head;
while (p) {
    process(p);
    p = p->next;
}
```

- the loop worksharing construct only works with loops **for**
- The **while** loops are not covered



Test2: linked lists the hard way

- Consider the program [linked.c](#)
 - Traverses a linked list computing a sequence of [Fibonacci numbers](#) at each node.
- Parallelize this program using constructs described so far.
- Once you have a correct program, optimize it.
- Code: `test2_Fibonacci/linked.c`



list traversal

- When OpenMP was first created, it focused on common use cases in HPC ... Fortran arrays processed over “regular” loops.
- Recursion and “pointer chasing” were removed from the Fortan focus.
- Hence, even a simple list traversal is exceedingly difficult with the original versions of OpenMP.



Test2-1: Linked lists omp25

```
while (p != NULL) {  
    p = p->next;  
    count++;  
}  
p = head;  
for(i=0; i<count; i++) {  
    parr[i] = p;  
    p = p->next;  
}  
#pragma omp parallel  
{  
    #pragma omp for  
    schedule(static,1)  
    for(i=0; i<count; i++)  
        processwork(parr[i]);  
}
```

Count number of items in the linked list

Copy pointer to each node into an array

Process nodes in parallel with a for loop

Code: test2_Fibonacci/solution/Linked_omp25.c



Linked lists

■ Using C++ STL

```
std::vector<node *> nodelist;
```

Count number of items in the linked list

```
for (p = head; p != NULL; p = p->next)  
    nodelist.push_back(p);
```

```
int j = (int)nodelist.size();
```

Count number of items in the linked list

```
#pragma omp parallel for schedule(static,1)
```

```
    for (int i = 0; i < j; ++i)  
        processwork(nodelist[i]);
```

Process nodes in parallel with a for loop



Conclusion

- We were able to parallelize the linked list traversal ...
 - but it was **ugly** and **required multiple passes** over the data.
- To move beyond its roots in the array based world of scientific computing, we needed to:
 - support **more general data structures**
 - support loops beyond basic 'for' loops.
- To do this, we have the **OpenMP 3.0**



references

- Tim Mattson, A “Hands-on” Introduction to OpenMP, Intel Corp.
- Mark Bull, Parallel Programming with OpenMP, EPCC, University of Edinburgh, UK.
- 迟学斌等，并行计算与实现技术，科学出版社。