

Introduction to OpenMP Part4. Advanced Topics

人工智能学院 缪青海 miaoqh@ucas.ac.cn



Content

- Getting started with OpenMP
- Core features of OpenMP
- Work with OpenMP
- Advanced topics of OpenMP



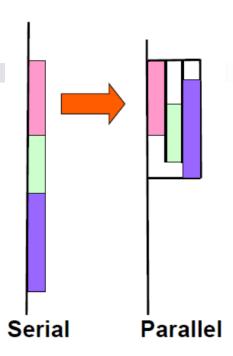
Advances of OpenMP

- Tasks for linked list
- Monte Carlo Calculations
- Parallel Patterns



OpenMP Tasks

- Tasks are independent units of work.
- Tasks are composed of:
 - □ **code** to execute
 - data environment
 - □ internal control variables (ICV)



- Threads perform the work of each task.
- The runtime system decides when tasks are executed
 - □ Tasks may be deferred
 - □ Tasks may be executed immediately



Task Definitions

- **■** Task construct
 - □ task directive plus structured block
- Task
 - the package of code and instructions for allocating data created when a thread encounters a task construct
- Task region
 - □ the dynamic sequence of instructions produced by the execution of a task by a thread



Task Construct – Explicit Tasks

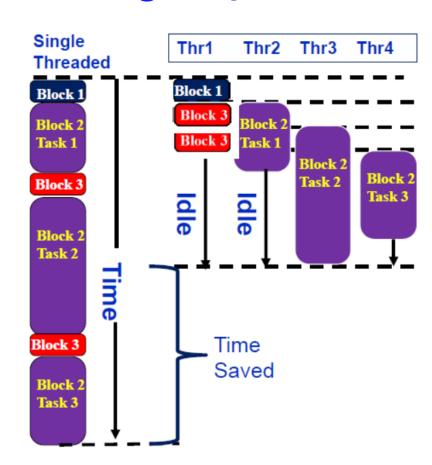
1. Create a team of #pragma omp parallel threads. #pragma omp single 2. One thread 3. The "single" thread executes the node * p = head; creates a task with its own **single** construct value for the pointer p while (p) { ... other threads #pragma omp task firstprivate(p) wait at the implied process(p); barrier at the end of p = p-next;the single construct 4. Threads waiting at the barrier execute tasks. Execution moves beyond the barrier once all the tasks are complete



Execution of tasks

Have potential to parallelize irregular patterns and

recursive function calls





When tasks completed?

- Tasks are guaranteed to be complete at thread barriers:
 - □ #pragma omp barrier
- or task barriers
 - ☐ #pragma omp taskwait

```
#pragma omp parallel

#pragma omp task

foo();

#pragma omp barrier

#pragma omp single

{

#pragma omp task

Cone bar task created here

bar();

bar task guaranteed to be completed here

completed here
```



- Fibonacci example:
 - ☐ The divide and conquer design pattern

```
int fib ( int n )
{
    int x,y;
    if ( n < 2 ) return n;
    #pragma omp task
        x = fib(n-1);
    #pragma omp task
        y = fib(n-2);
    #pragma omp taskwait
    return x+y
}</pre>

    n is private in both tasks
        x is a private variable
        y is a private variable
        y is a private variable
        What's wrong here?
```

A task's private variables are undefined outside the task



- Fibonacci example:
 - □ The divide and conquer design pattern

```
int fib ( int n )
{
    int x,y;
    if ( n < 2 ) return n;
    #pragma omp task shared (x)
        x = fib(n-1);
    #pragma omp task shared (y)
        y = fib(n-2);
    #pragma omp taskwait
    return x+y
}</pre>

    n is private in both tasks

    x & y are shared
    Good solution
    we need both values
    to compute the sum
```



List Traversal example:

```
List ml; //my_list
Element *e;

#pragma omp parallel

#pragma omp single
{
    for(e=ml->first;e;e=e->next)
        #pragma omp task
        process(e);
}

What's wrong here?

Possible data race!
Shared variable e updated by multiple tasks

tasks
```



List Traversal example:

```
List ml; //my_list
Element *e;

#pragma omp parallel

#pragma omp single

{
    for(e=ml->first;e;e=e->next)
        #pragma omp task firstprivate(e)
        process(e);
}
```



Test2-2: tasks in OpenMP

- Consider the program linked.c
 - ☐ Traverses a linked list computing a sequence of Fibonacci numbers at each node.
- Parallelize this program using tasks.
 - ☐ Compare new solution's complexity to the approach without tasks.

- Code:
 - □ test2_Fibonacci/solutions/linked_omp3_tasks.c



Advances of OpenMP

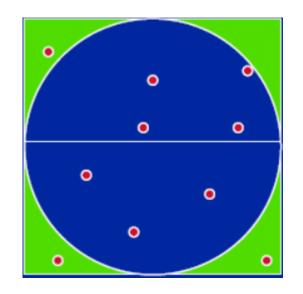
- Tasks for linked list
- **■** Monte Carlo Calculations
- Parallel Patterns



Monte Carlo Calculations



- Monte Carlo Method:
 - □ Using Random numbers to solve tough problems
 - □ Sample a problem domain to estimate areas, compute probabilities, find optimal values, etc.
- **Example:** Computing π with a digital dart board:



$$N = 10$$
 $\pi = 2.8$

$$N=100$$
 $\pi = 3.16$

$$N = 1000 \quad \pi = 3.148$$



Monte Carlo Calculations:



Chance of falling in circle is proportional to ratio of areas:

$$\square$$
 Ac = $\pi * r^2$

$$\square$$
 As = (2*r) * (2*r) = 4 * r²

$$\square$$
 P = Ac/As = π /4

- Compute π by randomly choosing points, count the fraction that falls in the circle, compute pi.
- Code: test3_Pi_mc/pi_mc.c 串行代码



Monte Carlo Calculations:

- 20
 - Parallel Programmers love Monte Carlo algorithms.
 - **■** Embarrassingly parallel:
 - ☐ the parallelism is so easy its embarrassing.
 - Add two lines and you have a parallel program.
 - Code: test3_Pi_mc/solutions/pi_mc_par.c 并行代码



Linear Congruential Generator (LCG)

LCG: Easy to write, cheap to compute, portable, OK quality.

```
random_next = (MULTIPLIER * random_last + ADDEND)% PMOD;
random_last = random_next;
```

- If you pick the multiplier and addend correctly, LCG has a period of PMOD.
- Picking good LCG parameters is complicated, so look it up:
 - \square MULTIPLIER = 1366
 - \square ADDEND = 150889
 - \square \square PMOD = 714025
- Code: test3_Pi_mc/solutions/random_seq_lcg.c

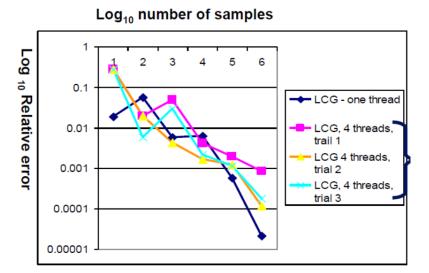
```
compile: gcc -fopenmp pi_mc_par.c random_seq_lcg.c
```



PI_MC with LCG generator

■ Problem:

- □ Run the same program the same way and get different answers!
- ☐ That is not acceptable!



Issue:

☐ the LCG generator is not threadsafe.



LCG code: threadsafe version

```
static long MULTIPLIER = 1366;
static long ADDEND = 150889;
static long PMOD = 714025;
                                 Construct (指令)
long random_last = 0;
#pragma omp threadprivate(random last)
double random ()
                            threadprivate不是针对某一个并
                            行块,而是作用于整个程序。
 long random next;
 random_next = (MULTIPLIER * random_last + ADDEND)% PMOD;
 random last = random next;
 return ((double)random next/(double)PMOD);
```



LCG code: threadsafe version

- #pragma omp threadprivate(random_last)
- random_last carries state between random number computations
- To make the generator threadsafe, make random_last threadprivate so each thread has its own copy.
- Code:
 - □ test3_Pi_mc/solutions/random_par.c
 compile: gcc –fopenmp pi mc par.c random_par.c



Data sharing: Threadprivate

- Makes global data private to a thread
 - ☐ C: File scope and static variables, static class members
- Different from making them PRIVATE 区别
 - with PRIVATE global variables are masked in a parallel region.
 - ☐ THREADPRIVATE preserves global scope within each Thread.
- Threadprivate variables can be initialized using COPYIN or at time of definition (using language defined initialization capabilities).



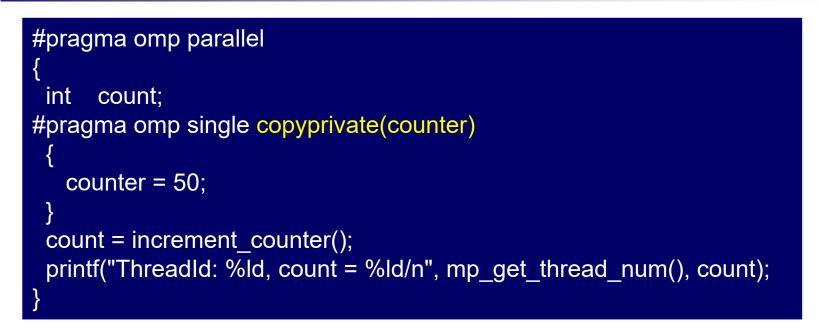
A threadprivate example

Use threadprivate to create a counter for each thread.

```
int counter = 0;
#pragma omp threadprivate(counter)
int increment_counter()
{
    counter++;
    return (counter);
}
```



A threadprivate example



Copyprivate():

☐ Broadcasts a value from the data environment of one implicit task to the data environments of the other implicit tasks belonging to the parallel region.



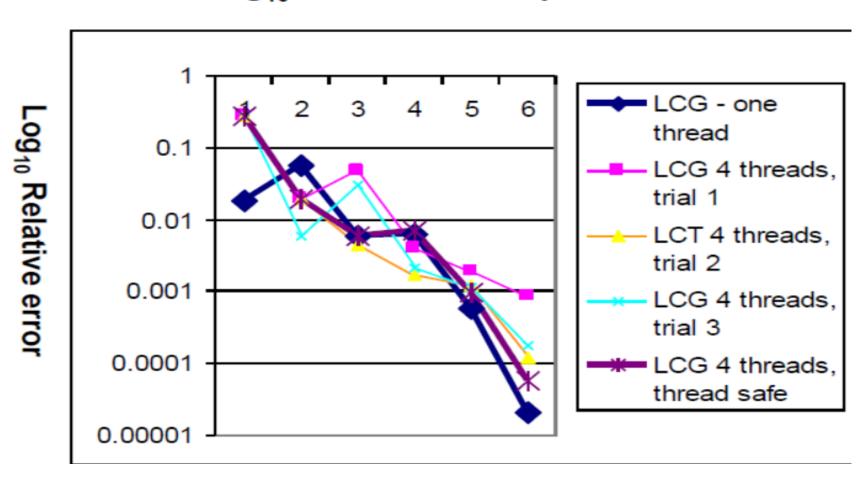
Thread safe random number generators

- test3_Pi_mc:
 - ☐ Thread safe version gives the same answer each time you run the program.
- But for large number of samples, its quality is lower than the one thread result!
- Why?



Thread safe random number generators

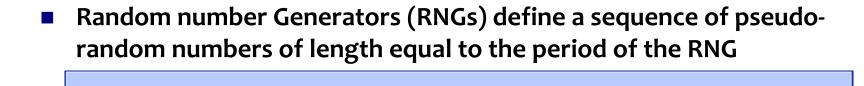
Log₁₀ number of samples





Thread 3

Pseudo Random Sequences



In a typical problem, you grab a subsequence of the RNG range

■ Grab arbitrary seeds and you may generate overlapping sequences

Thread 1

Thread 2

Overlapping sequences = over-sampling and bad statistics ... lower quality or even wrong answers!



Parallel random number generators



Solutions:

- ☐ Give each thread a separate, independent generator
- Have one thread generate all the numbers.
- Leapfrog ... deal out sequence values "round robin" as if dealing a deck of cards.
- □ Block method ... pick your seed so each threads gets a distinct contiguous block.

OR... use/buy a math library that does it right.



Advances of OpenMP

- Tasks for linked list
- Monte Carlo Calculations
- Parallel Patterns



Summary of OpenMP

- **Parallel Patterns:**
 - ☐ SPMD: Single Program Multiple Data
 - Loop parallelism
 - □ Divide and Conquer
 -



Parallel Patterns



■ SPMD:

- □ Run the same program on P processing elements;
- □ Use the rank ... an ID ranging from 0 to (P-1) ... to select between a set of tasks.
- □ most commonly used pattern in the history of parallel programming
- □ Example : computing Pi



SPMD

```
#include <omp.h>
void main (int argc, char *argv[])
    int i, pi=0.0, step, sum = 0.0;
    step = 1.0/(double) num steps;
    #pragma omp parallel firstprivate(sum) private(x, i)
      int id = omp_get_thread_num();
      int numprocs= omp_get_num_threads();
      int step1 = id *num_steps/numprocs;
      int stepN = (id+1)*num_steps/numprocs;
      if (stepN! = num_steps) stepN= num_steps;
      for (i=step1; i<stepN; i++){</pre>
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
      #pragma omp critical
          pi += sum *step;
```



Parallel Patterns



- Collections of tasks are defined as iterations of loops;
- Loop iterations are divided between a collection of processing elements.
- □ heavily used with data parallel design patterns;
- □ commonly used by OpenMP programmers.
- □ Example : computing Pi



LOOP

```
#include <omp.h>
      static long num_steps= 100000;
double step;
      #define NUM_THREADS 2
      void main()
          Int i; double x, pi, sum =0.0;
          step = 1.0/(double) num_steps;
          omp_set_num_threads(NUM_THREADS);
          #pragma omp parallel for private(x) reduction (+:sum)
              for (i=0; l < num steps; i++){
                x = (i+0.5)*step;
                sum += 4.0/(1.0+x*x);
          pi = sum * step;
```



Parallel Patterns

- Divide and Conquer:
 - □ Divide a problem into sub-problems;
 - □ Computing in parallel;
 - and recombine solutions of sub-problems into a global solution.
 - ☐ Example : computing Pi



```
#include <omp.h>
                                                                   Divide and
static long num steps= 100000000;
#define MIN_BLK 10000000
                                                                   Conquer
double pi_comp(int Nstart, int Nfinish, double step){
    Int i, iblk; double x, sum = 0.0, sum1, sum2;
    if (Nfinish-Nstart< MIN BLK){</pre>
                                                 int main (){
         for (i=Nstart; i< Nfinish; i++){</pre>
                                                   inti;double step, pi, sum;
                                                   step = 1.0/(double) num steps;
         x = (i+0.5)*step;
                                                   #pragma omp parallel {
         sum = sum + 4.0/(1.0+x*x);
                                                     #pragma omp single
                                                     sum=pi comp(o,num_steps,step);
    else{
                                                   pi = step * sum;
         iblk= Nfinish-Nstart;
         #pragma omp task shared(sum1)
             sum1 = pi_comp(Nstart, Nfinish-iblk/2, step);
         #pragma omp task shared(sum2)
             sum2 = pi comp(Nfinish-iblk/2, Nfinish, step);
         #pragma omp taskwait
             sum = sum1 + sum2;
    return sum;
```



Advances of OpenMP

- **■** More topics
 - ☐ Memory model
 - □ Teams
 - □ Distributed SIMD
 - □ Pairwise Synchronization
 -

https://www.openmp.org/specifications/

A separate County reference and for Fortran & also available



OpenMP API 3.1 C/C++ Directives An OperAtP executable directive applies to the sacreeling structured block or an OpenAP Construct. It structured block is a single statement or a compount statement with a single entry at the top and a single ent Parallel D.40 The parallel construct forms a team of threads and starts parallel execution. Fpragma one parallel (circus [[]circus [...] dischared dock nan Dreadchteprepressie) defeatiblemed I record frequirement (A) reduction operators into oce (2.5.1) The loop construct specifies that the berations of ops will be distributed among and executed by the recountering team of threads. Eprogram comp for [closes [[,]closes] ...] private (NC) frequience (in) insprience (in) Most common form refection/operators list) for par + by cheddeland, class size to the indicated about size as churics are scheduled the ran-solvel vor ICX

static terations are diskled into churic of size chank size. Chunks are assigned to threads in the then requests another churic until no churic remain golded: Each Coreal executes a churck of berations then requests another churic until no churic remain to be assigned. The churic stars start large and shrink

ser relational-op b

ser ++ (no)

auto. The decision regarding scheduling is delegated to the compiler and/or runtime setters. untine. The schedule and churic size are taken from

that are to be distributed among and executed by the

soundering team of threads.

Paragraph one section (descript class)

```
Personal sing saction?
Programme one section
```

and the second

The single construct specifies that the associated structured block is executed by only one of the threads is the team (not necessarily the master thread), in the centent of its implicit task

OpenMP 3.1 API C/C++ Syntax Quick Reference Card

(S.D.D) refers to sections in the OperAFF API Specification available at www.openmp.org.

OpenMP Application Program Interface (MPI) is a portable, scalable. OpenMP supports multi-platform shared-memory parallel

modelthat gless shared memory parallel programmers a simple programming to C/CHH and Fortram on all architectures, including and feature transfer for developing parallel applications for this platforms and Windows NT platforms.

```
tyregree cop chale (discre) [, histor) ... ]
```

platforms ranging from the desistop to the supersompates.

applete(kt)

Parallel Loop (2.6.1)

The parallel loop construct is a shortcut for specifying a sampled construct containing one or more associated

Opengree comp parallel for [classe] [, [classe] ...]

Are accepted by the panillel or for directives, except the nowalt clause, with identical meanings and restActions.

Simple Parallel Loop Example

The following example demonstrates how to parallelize a simple toop using the parallel toop

well simplected at Back for Back from

bysages may presided the for their transition J^{n} is to potentially default $^{n}/$ h(h)=(a(h)+a(h+1)) / 2.4τ

The parallel sections construct is a chartcut for specifying a parellel construct containing one sections construct an to other statements.

Pyragma omp parallel sections (classes (_)classes [__)

Pyragma omp saction) Pyragna one section

Any of the clauses accepted by the parallel or sections directives, except the nowalt clause, with identical

The task construct defines an explicit task. The data eretroment of the task is prested according to the data-sharing attribute classes on the task sometrust

and any defaults that apply. Program con task (discor)(,)drass)...]

Expiremental for facility expression?

debuilt(classed | none)

The taskyteld construct specifies that the current task can be suppressed in favor of execution of a different task.

is executed by the master thread of the team. There is no implied barrier either on entry to, or exit from, the

Gritical [2.8.2] The critical corell ruct redricts execution of the associated structured block to a single thread at a time.

The banter construct specifies an explicit banker at the point at which the construct appears.

Epingma omp barrier

The Seised construct specifies a wait on the completion of child tasks of the current task

Atomic [2.8.5]

The attents constitut ensures that a specific storage ration is updated atomically, rather than expessing it to the appointing of multiple, simplements writing threads.

Februari I	eprodución i				
-	170				
erio .	411000				
opideor had proof	7	a bitage variety	ere likepape		
anjan .		ena Magnage;	14.00		
and discretized block may be one of the following frame					

Purb [2.8.6]

The flash construct executes the OperAP flash operation, which makes a thread's temporary view of memory consistent with memory, and entirese an order on the memory operations of the variables.

Ordered [2.8.7] The ordered conduct specifies a structured block to a

Threadprison [2,9,2] The Greedprison director specifies that vertables are

A commo reparated list of file-scope, namespace scope, or static block-scope variables that do not have incomplete types.

Page 1

Foregrap omp technick

Marter [2.8.1] The master construct specifies a structured block that

fpragma omp mader

fungma omp chical (name)

bolwet [J.J.4]

Fprogram omp technolit

foregraphic and stank lead (write) update (quotien)

foregree one stank against

where expression start may be one of the following forms:

der b.	eprodu			
-	170			19
rise .	a magan			9
alder al passi	4	a littley mappy	and dispression	=
phore	*****	energy (14.00	**

(ichteuprespo school) frequentling-expc) (coadingroup) vos)

improgram that will be executed to the order of the large fleviations. This requestratives and orders the code-will to an ordered region-white allowing code outside the region.

foregree one ordered

replicated, with each thread having its own copy.

Foregrap omp (freedprivate (fr.))

OpenMP API 3.1 C/C++

Runtime Library Routines

Execution Environment Routines (3.2) becation englowment realizes affect

and monitor threads, processors, and the sarafai envisorment

Misca the number of threads used for subsequent personal revolves that do not

beturns the number of threads in the Carrent Seam.

int congress, many threads both; som sunder of threads than parellel construct without a rare. Streads

fetures the ID of the encountering thread where 10 ranges from serv to the size of

int amp get, man procedured;

the team release 1.

teturns the number of processors available to the program.

Dieta Types For Bustime Ultrary Routines

any jack t. Represents a simple lack ong_lest_look_2 Represents a

one what it bepresents a schedule

being to people vall. Returns tree if the call to the routine is enclosed by an active parallel region;

trables or disables dynamic adjustment of the number of threads available by setting the value of the abovar ICX.

Nations the value of the dynesor ICV. determining whether dynamic adjustment of the number of threads is enabled or

will any set nested by reserving tradies or displies nested sarabelors, by setting the sent wor ICV.

Returns the value of the nest our CX; which determines if rented parallelon is enabled or disabled.

widomp or wheldelmp abell that

Affects the schedule that is applied when northre is used as schedule bind, by setting the value of the sur-schedule FOV.

and home of state, denumb, guided, auto, or an implementation-defined schedule, time loop-construct (2.6.5) for descriptions.

and one pet abeliand and modified

Returns the value of our-sphed-year ICV, which is the schedule applied when surtine schedule is used.

See Aired above.

the original them.

become get thread Britishilly

fatures the value of the thread-destroyor CV, which is the madesum number of

int may learly Units the number of nested active parellel region, by setting more of the dress war

transport, narrathe, braid with National Devalue of managing law investigate

number of restal active smallel regions. (disclosed top general

Returns the number of nested parallel regions enclosing the task that contains

intump_get_encedor_threat_nam(Natures, for a given nested less of the cornect thread, the thread number of the ancestor or the current thread.

Natures, for a given needed level of the cornect Downs, the size of the thread team

to which the ancestor or the correct firear bit amp get author involved) Returns the number of nested, activ parallel regions enclosing the task that contains the call.

bit sing in final sold; Natures true if the routine is executed in a final or included task region; otherwise, it

Lock Routines [3,3]

Lock routines support synchronization with-

rold anguist, judgeng, judge Nody. with many last ment back!

that I had gradied yetter age blos and the food bear dead

These routines ensure that the OpenMF took is untritlettend.

with any one builting, builty finish ong net jod j find)

an OperAP lock.

will any west indices; but I had

and and parameters for the These routines provide a means of smelting an Openitiff lock.

belong bed Jedjong Jock ! Not?

These routines alternot to set an OperAFF tack but do not suspend mercition of the task merciting the routine.

Tirring motines support a portable wall

double only get witnessed.

docks one per whitever) facture the precision of the firmer used by using gar, without

Clauses

articular direction is described with the traction. Hind clauses accept a some aparated list of list have, all list have meaning in a clause must be obtion.

uta Sharing Attribute-Clauses (3:8.3) beta-sharing attribute discess apply only

to variables whose names are stoble in the construct on which the clause accesss. Controls the default data-sharing attributes of variables that are referenced to a parallel

Declares one or more list literat to be chared by taxis: generated by a paraflehor task

private (in)

Declares one or more flot bens to be private to a bein.

transport

Declares one or more list bens to be private to a task, and institution each of there with The value that the corresponding original bern has when the constituct to encounter individual list

bedans one or more list being to be

after the end of the region.

or backe to an implicit task, and causes the

presponding original flees to be updated

Data Copying Clauses [2.9.4]

These classes apport the copying or data values from private or Evenigor variables on one implicit task or thread to the corresponding variables on other

Declares accumulation into the list have

using the indicated associative operator.

Species by reducing Difficulties return.

nin farest nation is reliable by ten tree!

currelation occurs into a private copy for its list bear which is then combined with

Copies the value of the master thread's threely hate variable to the threely hate variable of each other member of the team essecting the panelse region.

Broadcasts a write from the data environment of one implicit task to the data environments of the other triplicit balls

Environment Variables

traditionment sortables are described section [4] of the KP specification Environment sertable names are upper date, and the selves sedgred to their are case insensitive and may have leading

and trading white space. OMF SOMOUTH Symichania late the considering SV for the demands, golded, or solls, claim! It a

time integer that specifies shork size. OMF NUM THROUGH BUT Sets the otherwisear ICV for the number of threads to use for parellel regions.

OMP_DYNAMIC downs tats the dynesic ICV for the dynamic adjustment of threath to use for paradici regions. Valid relies for specimic are time

OMP PROC MIND NOW

has the value of the global bland our KY. The value of this environment variable road by true or false.

lefs the nearwor ICV to enable or to shadde nested parallellars, Valid values.

for mested are true or false. OMP_DINCHES CONT. (N) (I) the size of the stack for threads created by the OperAt/Proplementation, size is a office Integer that specifies stack size.

If unit is not specified, size to measured to OMF_MAT_POUCY pully Sets the work-polity-voir ICV that controls the destreal behavior of working threads. fallet values for polity are ACTMS

DAF MAI ACTIVE LEVEL levels lets the our active inveloper ICE that certain the madrium number of nested

OMP, THROUGH DATE THAT Sets the shread-simb our ICV that sonthisk the madenum number of threads perficulting in the OperAFF program.

provide & 2011 Deput AP South Residence Season Season, Personal resource of South Season and property The materials granted provided the OpenSP Architecture Series Scarie engagin and officed this description assess. Solidon is given that copying is by perceivalent of the OpenSP offices on Federal Season. Persistent or published on Season are one or or or of the OpenSP of Season Season. questifiair tions mod automataige the copyright by deplaying the following darkmann: "Operfoll" is a realizant of the OperAT Architecture Server Source, Neutron of the product publishation may been less realizant from the Operfoll anguage, Septimber Program interface (publishation).





OpenMP Application Programming Interface

Examples

- Supplementary Source Code GitHub Repository
- OpenMP 5.0 Examples Nov 2019
- OpenMP 5.0 Examples Discussion Forum



references

- https://www.openmp.org
- Tim Mattson, A "Hands-on" Introduction to OpenMP, Intel Corp.
- Mark Bull, Parallel Programming with OpenMP, EPCC, University of Edinburgh, UK.
- 迟学斌等,并行计算与实现技术,科学出版社。