



中国科学院大学
University of Chinese Academy of Sciences

Introduction to MPI

Part2. *Intermediate Topics* (A)

人工智能学院
缪青海

miaoqh@ucas.ac.cn



Content

- Non-blocking communication
- Group (collective) communication
- MPI Datatypes
- Virtual Topology



Mpi buffer

■ MPI定义三种缓冲区buffer:

□ 应用缓冲区a_buffer

- 将要发送/接收数据的地址空间，在消息格式中定义；

□ 系统缓冲区s_buffer

- MPI环境为通信所准备的存储空间，a_buffer在此出入复制；

□ 用户定义缓冲区u_buffer

- 用户在使用某些api时在程序中显式申请的用于通信的空间，注册到MPI环境中；




Blocking Communication

- In blocking communication:
 - **MPI_SEND** does not return until buffer is empty (available for reuse)
 - **MPI_RECV** does not return until buffer is full (available for use)
- A process sending data will be blocked until data in the send buffer is emptied.
- A process receiving data will be blocked until the receive buffer is filled.
- Exact completion semantics of communication generally depends on the **message size** and the system **buffer size**.




Blocking Communication


标准阻塞通信的状态:



立即发送



消息小于系统缓冲区:
缓存并返回
消息大于系统缓冲区:
等待接收



立即发送



Blocking Communication

阻塞通信的类型:

类型	发送	接收	注释
标准模式	MPI_Send	MPI_Recv MPI_Irecv MPI_Recv_Init	
缓冲模式	MPI_Bsend		wait
就绪模式	MPI_Rsend		test
同步模式	MPI_Ssend		



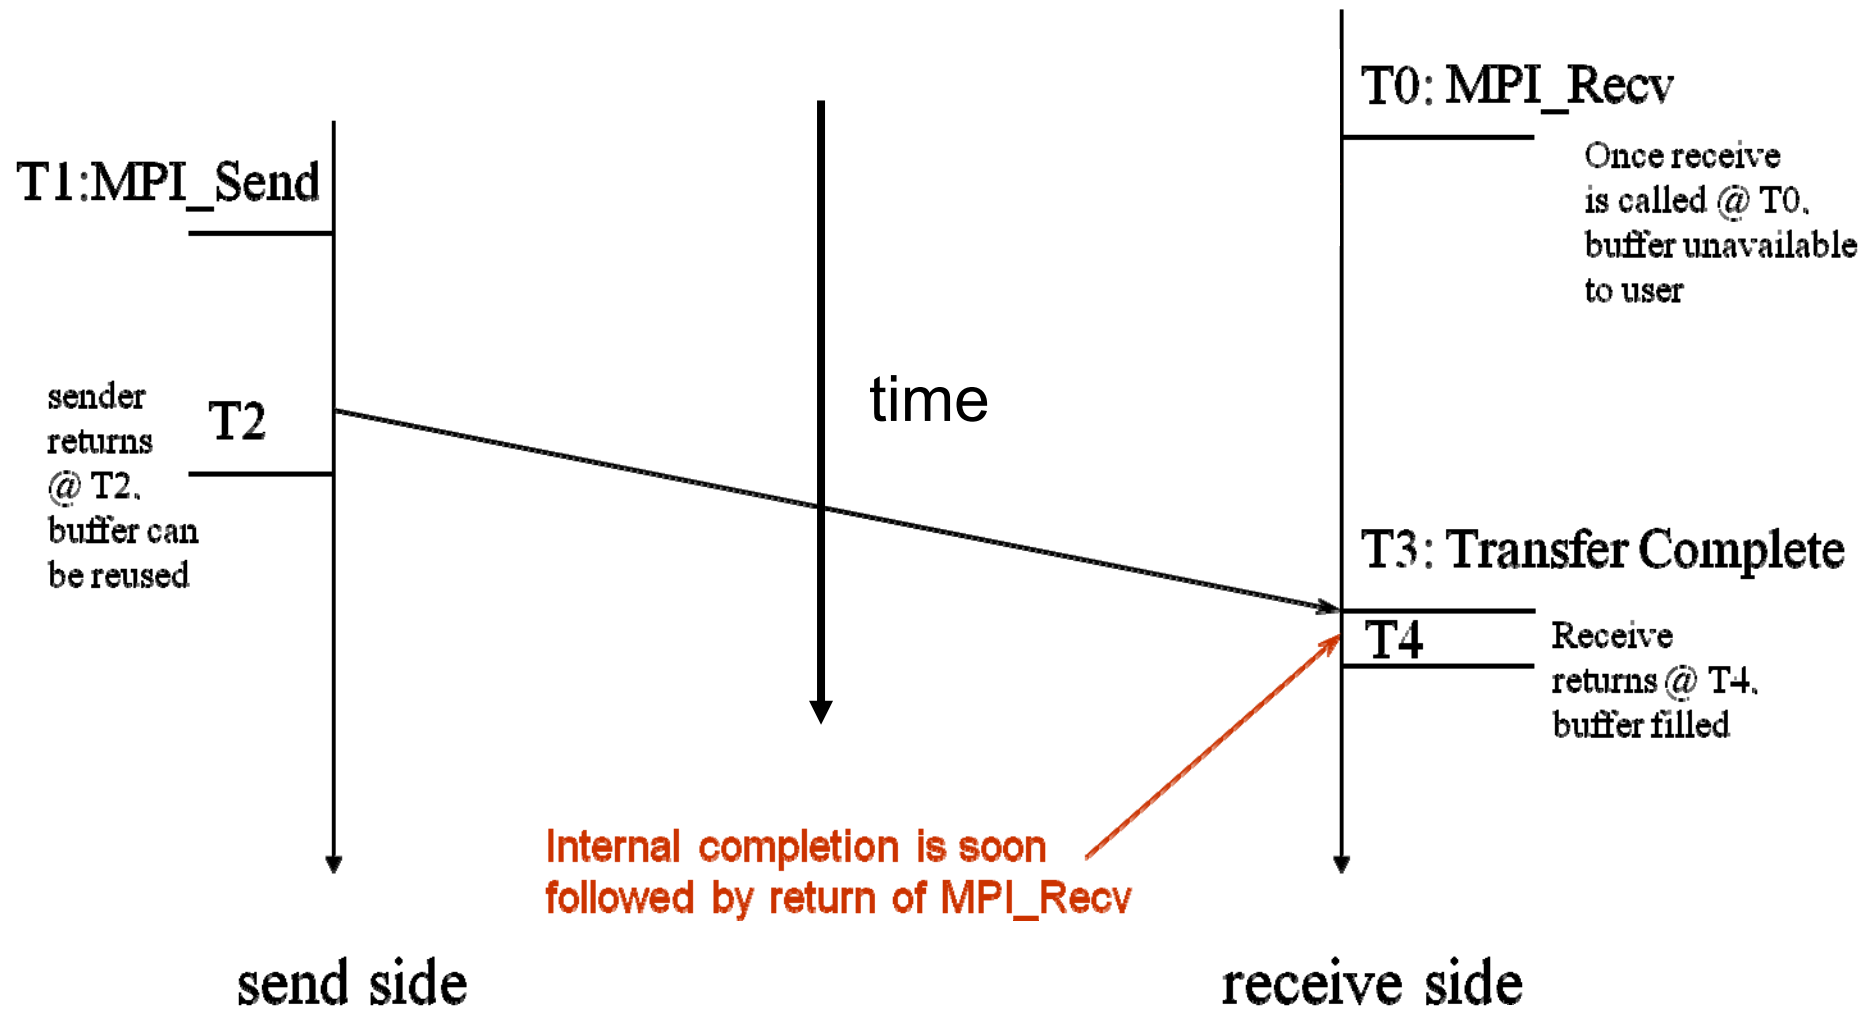
Blocking Communication

- In blocking communication.
 - **MPI_SEND** does not return until buffer is empty (available for reuse)
 - **MPI_RECV** does not return until buffer is full (available for use)
- Blocking communication is simple to use but can be prone to **deadlocks**

```
if (rank == 0) {  
    MPI_SEND(..to rank 1..)   
    MPI_RECV(..from rank 1..)   
    Usually deadlocks →  
else if (rank == 1) {  
    MPI_SEND(..to rank 0..) ← reverse send/recv  
    MPI_RECV(..from rank 0..)   
}
```



Blocking Send-Receive Diagram





Blocking vs. Non-blocking

- **MPI_SEND/MPI_RECV** are blocking communication calls
 - Return of the routine implies completion
 - When these calls return the memory locations used in the message transfer can be safely accessed for reuse
 - For “send” completion implies variable sent can be reused/modified
 - Modifications will not affect data intended for the receiver
 - For “receive” variable received can be read
- **MPI_ISEND/MPI_IRECV** are non-blocking variants
 - Routine returns immediately
 - – completion has to be separately tested for
 - These are primarily used to **overlap computation and communication** to improve performance



Non-Blocking Communication

- Non-blocking (asynchronous) operations return (immediately) “request handles” that can be waited on and queried
 - `MPI_ISEND(buf, count, datatype, dest, tag, comm, request)`
 - `MPI_Irecv(buf, count, datatype, src, tag, comm, request)`
 - `MPI_WAIT(request, status)`
- Anywhere you use **MPI_SEND** or **MPI_RECV**, you can use the pair of **MPI_ISEND/MPI_WAIT** or **MPI_Irecv/MPI_WAIT**
- One can also test without waiting using **MPI_TEST**
 - `MPI_TEST(request, flag, status)`




Multiple Completions

- It is sometimes desirable to wait on multiple requests:
 - `MPI_Waitall(count, array_of_requests, array_of_statuses)`
 - `MPI_Waitany(count, array_of_requests, &index, &status)`
 - `MPI_Waitsome(incount, array_of_requests, outcount, array_of_indices, array_of_statuses)`
- There are corresponding versions of `TEST` for each of these



Non-Blocking Send-Receive Diagram

High Performance Implementations
Offer Low Overhead for Non-blocking Calls





Message Completion and Buffering

- A send has completed when the user supplied buffer can be reused

```
*buf =3;  
MPI_Send(buf, 1, MPI_INT ...)  
*buf = 4; /* OK, receiver will  
always receive 3 */
```

```
*buf =3;  
MPI_Isend(buf, 1, MPI_INT ...)  
*buf = 4; /*Not certain if receiver  
gets 3 or 4 or anything else */  
MPI_Wait(...);
```

- Just because the send completes does not mean that the receive has completed
 - Message may be buffered by the system
 - Message may still be in transit



A Non-Blocking example

```
int main(int argc, char ** argv)
{
    [...snip...]
    if (rank == 0) {
        for (i=0; i< 100; i++) {
            /* Compute each data element and send it out */
            data[i] = compute(i);
            MPI_Isend(&data[i], 1, MPI_INT, 1, 0, MPI_COMM_WORLD,
                    &request[i]);
        }
        MPI_Waitall(100, request, MPI_STATUSES_IGNORE)
    }
    else {
        for (i = 0; i < 100; i++)
            MPI_Recv(&data[i], 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
                    MPI_STATUS_IGNORE);
    }
    [...snip...]
}
```



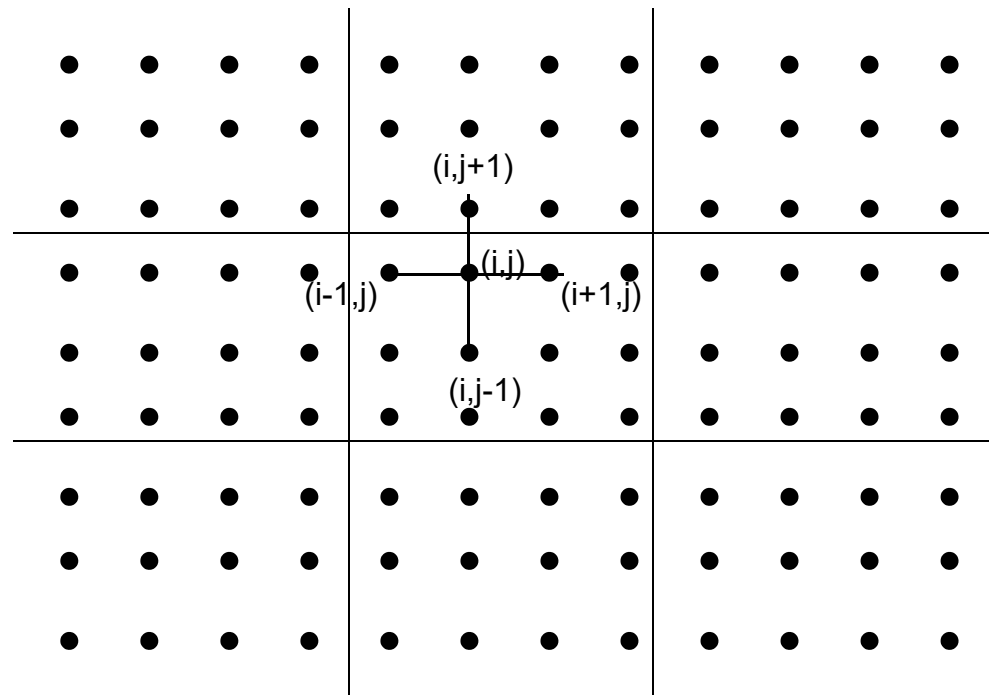
Code Example - 2

- `code/test2/mpi_helloNBsend.c`
- `code/test2/mpi_send_recv_nonblocking.c`

- `mpicc mpi_xxx.c -o test`
- `mpiexec -n 2 ./test`



2D Poisson Problem






Regular Mesh Algorithms

- Many scientific applications involve the solution of **partial differential equations (PDEs)**
- Many algorithms for approximating the solution of PDEs rely on forming a set of difference equations
 - **Finite difference, finite elements, finite volume**
- The exact form of the differential equations depends on the particular method
 - From the point of view of parallel programming for these algorithms, the operations are the same
- **Five-point stencil** is a popular approximation solution




Necessary Data Transfers






Necessary Data Transfers






Necessary Data Transfers




Necessary Data Transfers

- Provide access to remote data through a *halo* exchange (5 point stencil)



Necessary Data Transfers

- Provide access to remote data through a *halo* exchange (5 point stencil)






Understanding Performance:

■ Unexpected Hot Spots

- Basic performance analysis looks at two-party exchanges
- Real applications involve many simultaneous communications
- Performance problems can arise even in common grid exchange patterns
- Message passing illustrates problems present even in shared memory
 - Blocking operations may cause unavoidable memory stalls

Mesh Exchange

- Exchange data on a mesh





Sample Code

```
for (i = 0; i < n_neighbors; i++) {  
    MPI_Send(edge, len, MPI_DOUBLE, nbr[i], tag, comm);  
}  
for (i = 0; i < n_neighbors; i++) {  
    MPI_Recv(edge, len, MPI_DOUBLE, nbr[i], tag, comm, status);  
}
```

- What is wrong with this code?



Deadlocks!

- All of the sends may block, waiting for a matching receive (will for large enough messages)

- The variation of
if (**has up nbr**)

```
MPI_Recv( ... up ... )
```

...

```
if (has down nbr)
```

```
MPI_Send( ... down ... )
```

sequentializes (all except the bottom process blocks)




Fix 1: Use Irecv

```
for (i = 0; i < n_neighbors; i++) {  
    MPI_Irecv(edge, len, MPI_DOUBLE, nbr[i], tag,  
              comm, requests[i]);  
}  
for (i = 0; i < n_neighbors; i++) {  
    MPI_Send(edge, len, MPI_DOUBLE, nbr[i], tag, comm);  
}  
MPI_Waitall(n_neighbors, requests, statuses);
```

- Does not perform well in practice. Why?

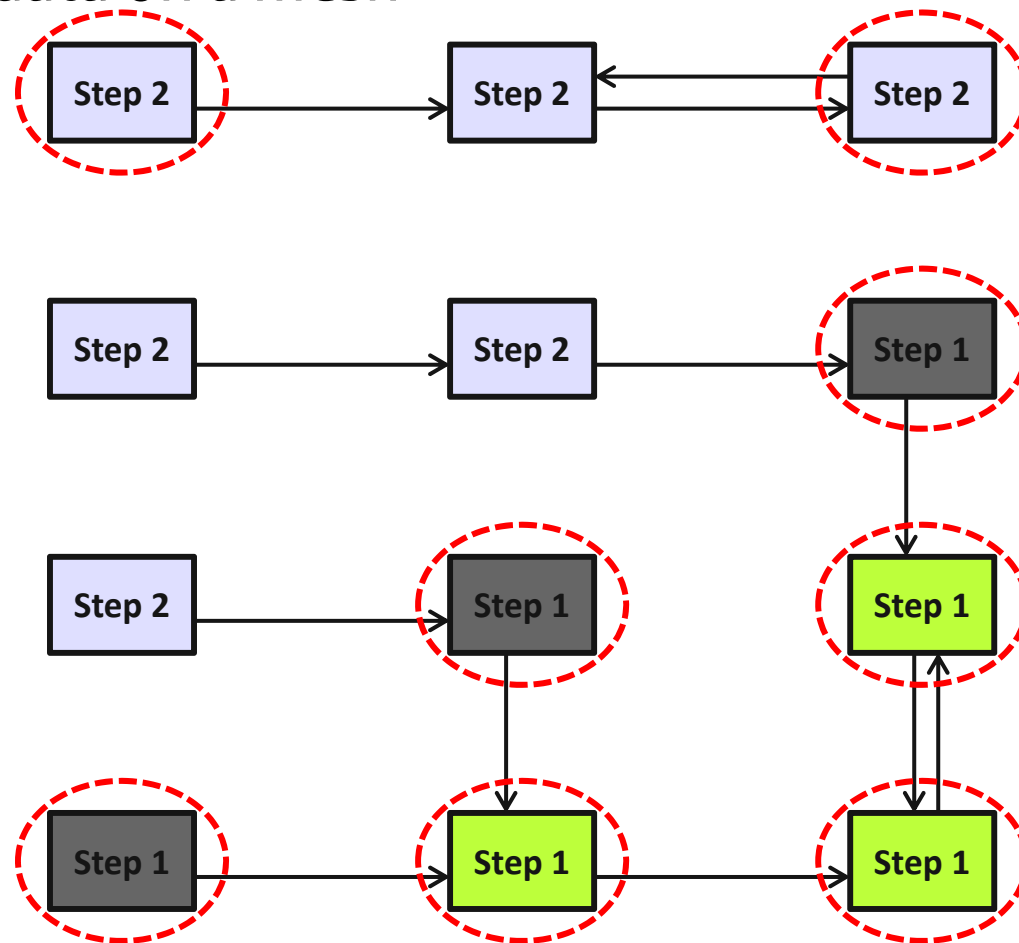
Mesh Exchange

- Exchange data on a mesh



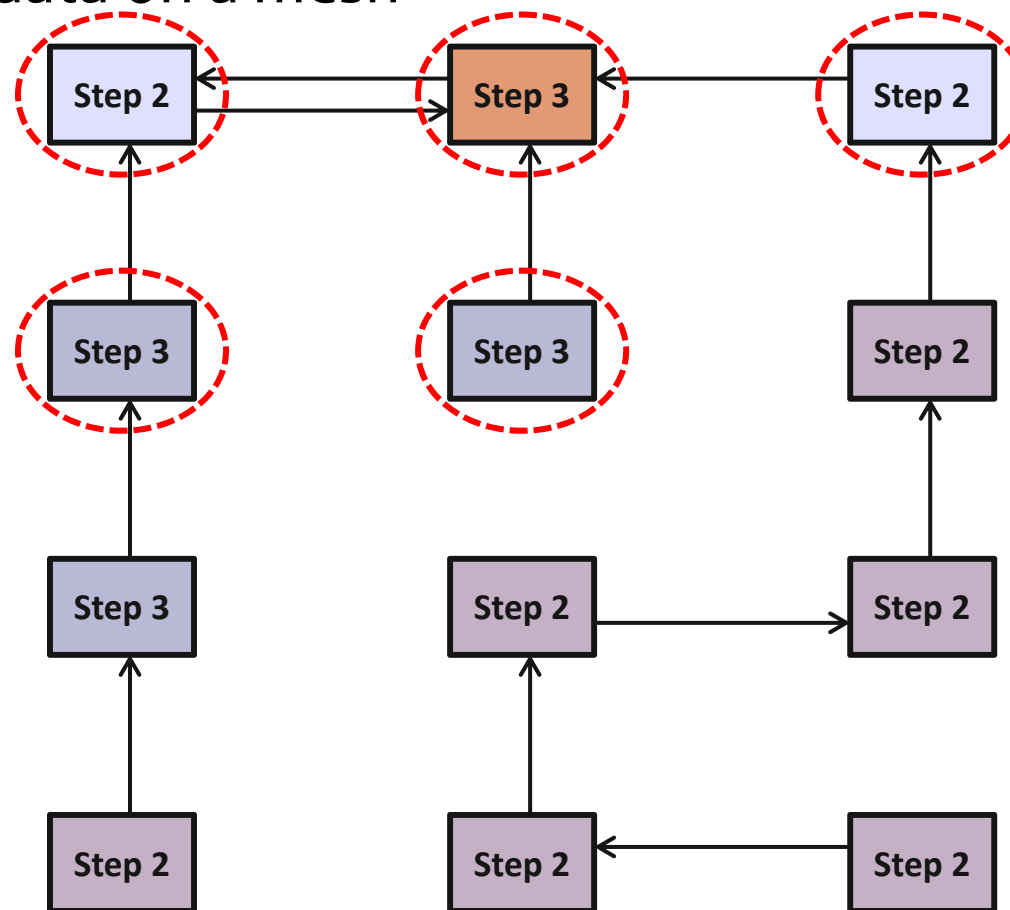
Mesh Exchange

- Exchange data on a mesh



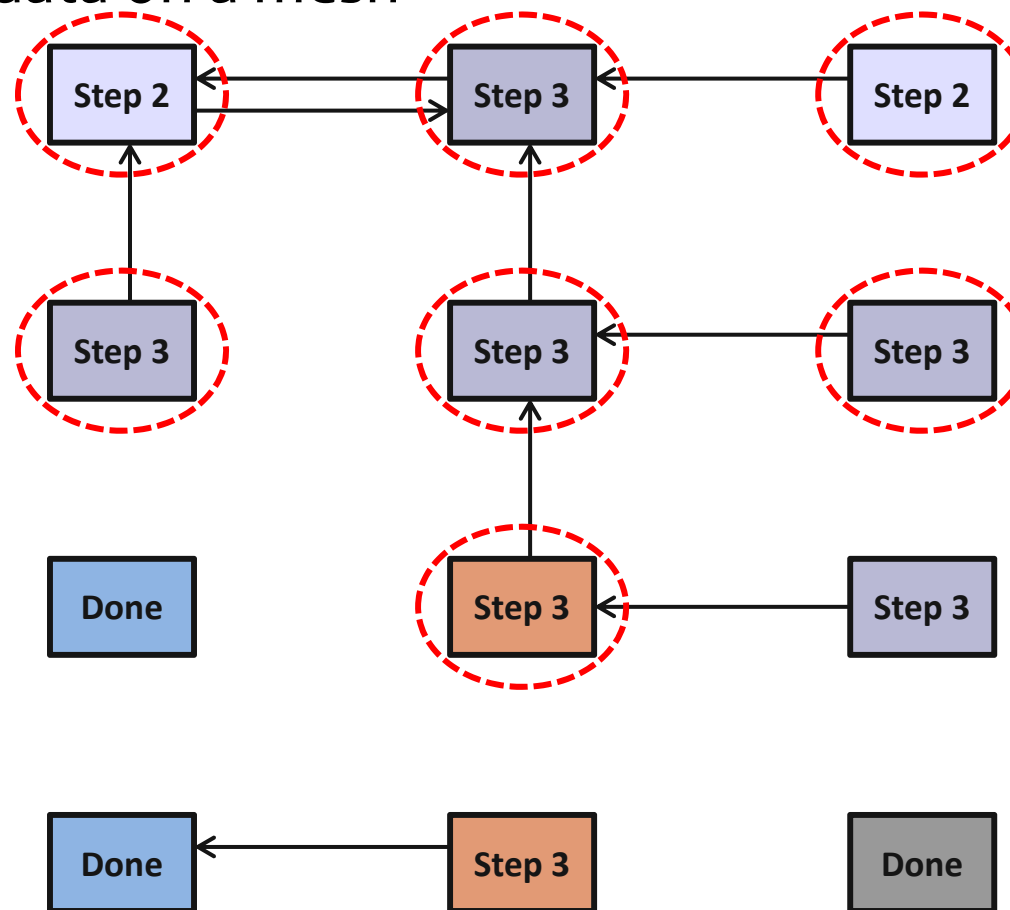
Mesh Exchange

- Exchange data on a mesh




Mesh Exchange

- Exchange data on a mesh



Mesh Exchange


- Exchange data on a mesh





Mesh Exchange

- Exchange data on a mesh



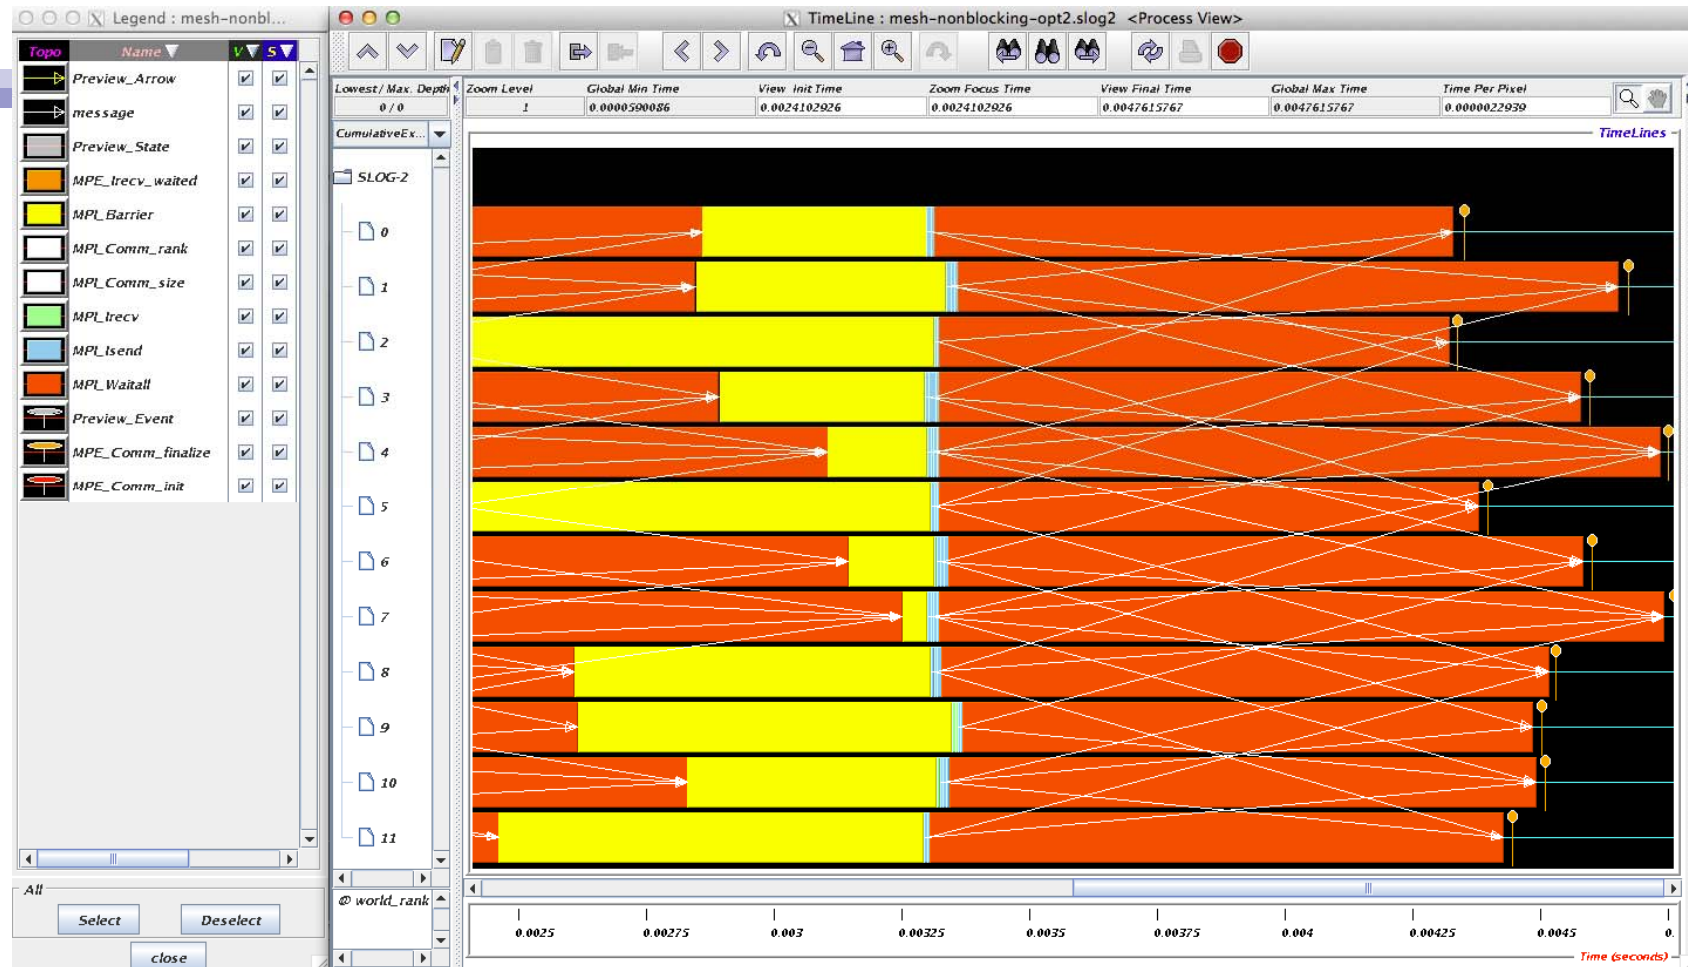


Fix 2: Use Irecv and Irecv

```
for (i = 0; i < n_neighbors; i++) {  
    MPI_Irecv(edge, len, MPI_DOUBLE, nbr[i], tag,  
              comm, requests[i]);  
}  
for (i = 0; i < n_neighbors; i++) {  
    MPI_Isend(edge, len, MPI_DOUBLE, nbr[i], tag, comm,  
             requests[n_neighbors + i]);  
}  
MPI_Waitall(2 * n_neighbors, requests, statuses);
```



Timeline from IB Cluster



Note processes 4 and 7 are the only interior processors; these perform more communication than the other processors



Lesson: Defer Synchronization

- Send-receive accomplishes two things:
 - Data transfer
 - Synchronization
- In many cases, there is more synchronization than required
- Use non-blocking operations and `MPI_waitall` to defer synchronization
- Tools can help out with identifying performance issues:
 - `MPE`, `Tau` and `HPCToolkit` are popular profiling tools
 - `Jumpshot` tool uses their datasets to show performance problems graphically
 - Display message queue state using `Totalview`



Content

- Non-blocking communication
- Group (collective) communication
- MPI Datatypes
- Virtual Topology



Collective Operations in MPI

- Collective operations are called by all processes in a communicator:
 - **MPI_BCAST**
 - distributes data from one process (the root) to all others in a communicator.
 - **MPI_REDUCE**
 - combines data from all processes in the communicator and returns it to one process.
- In many numerical algorithms:
 - **SEND/RECV** can be replaced by **BCAST/REDUCE**,
 - improving both simplicity and efficiency.



MPI Collective Communication

- Communication and computation is coordinated among a group of processes in a communicator.
 - Tags are not used;
- Three classes of operations:
 - **synchronization**
 - **data movement**
 - **collective computation**
- Non-blocking collective operations in MPI-3



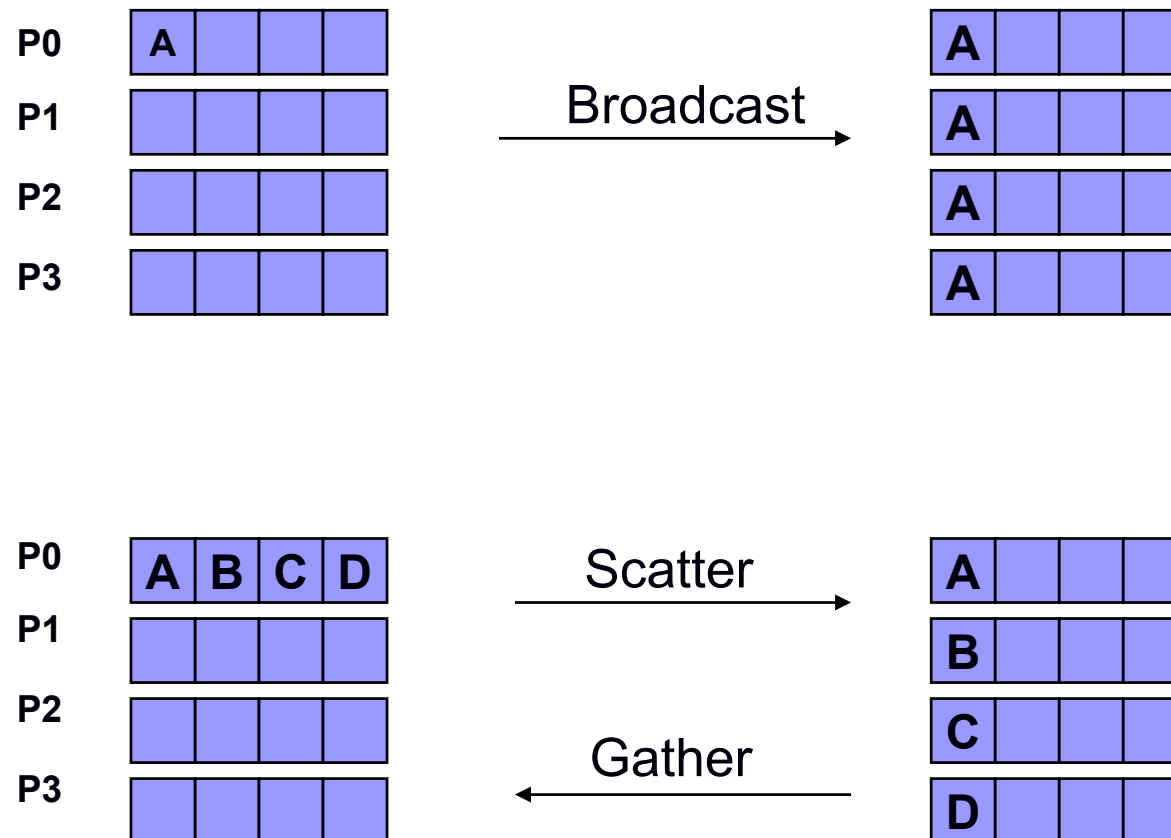
Synchronization

■ `MPI_BARRIER(comm)`

- ☐ Blocks until all processes in the group of the communicator `comm` call it.
- ☐ A process cannot get out of the barrier until all other processes have reached barrier.




Collective Data Movement







More Collective Data Movement





Collective Computation





MPI Collective Routines

- Many Routines:
 - `MPI_ALLGATHER`, `MPI_ALLGATHERV`, `MPI_ALLREDUCE`, `MPI_ALLTOALL`, `MPI_ALLTOALLV`, `MPI_BCAST`, `MPI_GATHER`, `MPI_GATHERV`, `MPI_REDUCE`, `MPI_REDUCESCATTER`, `MPI_SCAN`, `MPI_SCATTER`, `MPI_SCATTERV`
- “All”
 - versions deliver results to all participating processes
- “V”
 - versions (stands for vector) allow the chunks to have different sizes
- `MPI_ALLREDUCE`, `MPI_REDUCE`, `MPI_REDUCESCATTER`, and `MPI_SCAN` take both built-in and user-defined combiner functions



MPI Built-in Collective Computation

■ MPI_MAX	Maximum
■ MPI_MIN	Minimum
■ MPI_PROD	Product
■ MPI_SUM	Sum
■ MPI_LAND	Logical and
■ MPI_LOR	Logical or
■ MPI_LXOR	Logical exclusive or
■ MPI_BAND	Bitwise and
■ MPI_BOR	Bitwise or
■ MPI_BXOR	Bitwise exclusive or
■ MPI_MAXLOC	Maximum and location
■ MPI_MINLOC	Minimum and location



Defining your own Collective Operations

- Create your own collective computations with:

```
MPI_OP_CREATE(user_fn, commutes, &op);
```

```
MPI_OP_FREE(&op);
```

```
user_fn(invec, inoutvec, len, datatype);
```

- The user function should perform:

```
inoutvec[i] = invec[i] op inoutvec[i];
```

for i from 0 to len-1


- The user function:

- ☐ can be non-commutative;
- ☐ but must be associative.



Example: Calculating Pi

- Calculating pi via numerical integration
 - Divide interval up into subintervals
 - Assign subintervals to processes
 - Each process calculates partial sum
 - Add all the partial sums together to get pi



1. Width of each segment (w) will be $1/n$
2. Distance ($d(i)$) of segment “ i ” from the origin will be “ $i * w$ ”
3. Height of segment “ i ” will be $\sqrt{1 - [d(i)]^2}$



Example: PI in C

```
#include <mpi.h>
#include <math.h>
int main(int argc, char *argv[])
{
    [...snip...]
    /* Tell all processes, the number of segments you want */
    MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
    w    = 1.0 / (double) n;
    mypi = 0.0;
    for (i = rank + 1; i <= n; i += size)
        mypi += w * sqrt(1 - (((double) i / n) * ((double) i / n)));
    MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
               MPI_COMM_WORLD);
    if (rank == 0)
        printf("pi is approximately %.16f, Error is %.16f\n", 4 *
pi, fabs((4 * pi) - PI25DT));
    [...snip...]
}
```




Code Example - 3

- `code/test3/`
 - `cpi.c`
 - `mpi_group_scatter.c`
 - `mpi_group_reduce_pi.c`
 - `mpi_group_communicater.c`



Content

- Non-blocking communication
- Group (collective) communication
- [MPI Datatypes](#)
- Virtual Topology