



中国科学院大学
University of Chinese Academy of Sciences

Introduction to MPI

Part2. *Intermediate Topics* (B)

人工智能学院
缪青海

miaoqh@ucas.ac.cn



Content

- Non-blocking communication
- Group (collective) communication
- [MPI Datatypes](#)
- Virtual Topology



Simple/Predefined Datatypes

- Equivalents exist for all C, C++ and Fortran native datatypes:
 - C int → MPI_INT;
 - C float → MPI_FLOAT
 - C double → MPI_DOUBLE
 - C uint32_t → MPI_UINT32_T




Complex Datatypes

- For more **complex** or **user-created** datatypes, MPI provides routines to represent them as well:
 - ☐ Contiguous
 - ☐ Vector/Hvector
 - ☐ Indexed/Indexed_block/Hindexed/Hindexed_block
 - ☐ Struct
 - ☐ Some convenience types (e.g., subarray)
 - ☐


Necessary Data Transfers

- Provide access to remote data through a *halo* exchange (5 point stencil)



Necessary Data Transfers


- Provide access to remote data through a *halo* exchange (5 point stencil)





The Local Data Structure

- Each process has its local “patch” of the global array
 - “bx” and “by” are the sizes of the local array
 - Always allocate a halo around the patch
 - Array allocated of size $(bx+2) \times (by+2)$





Introduction to Datatypes in MPI

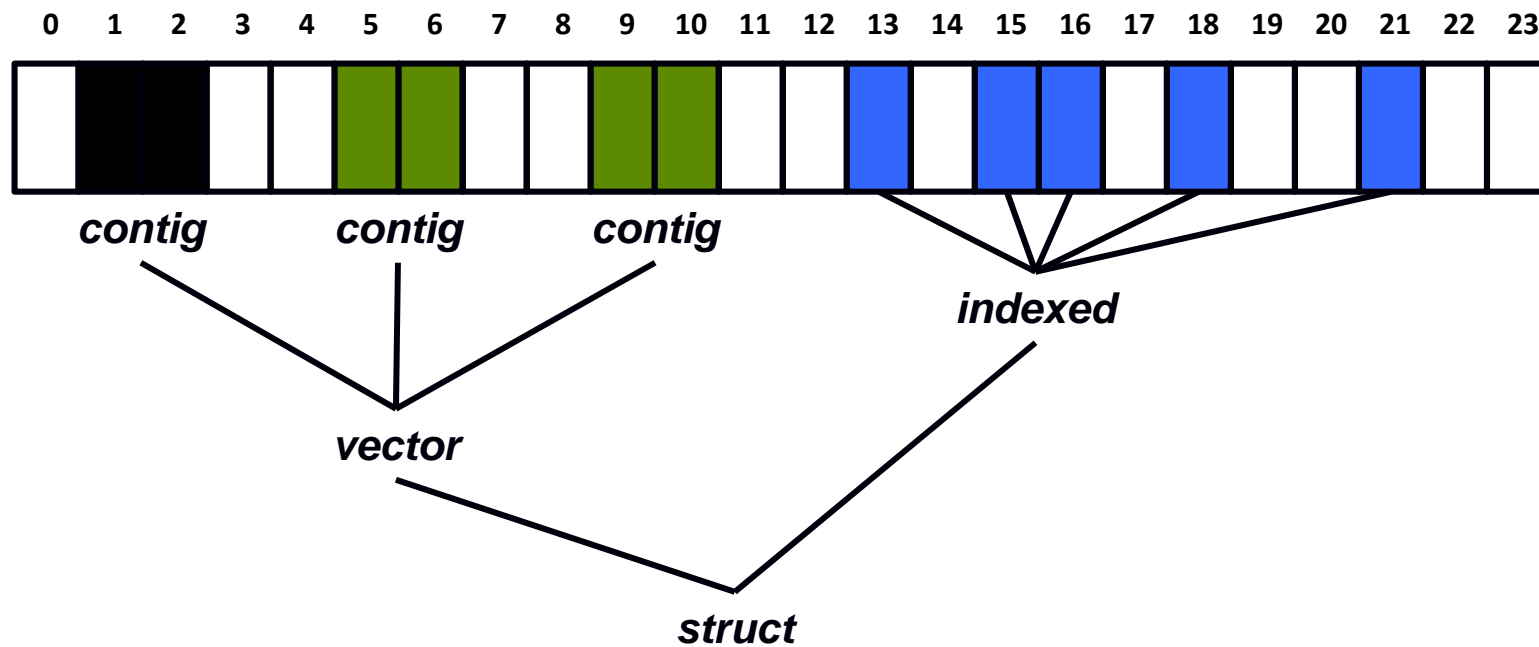
- **Datatypes** allow to (de)serialize **arbitrary** data layouts into a message stream
 - Networks provide serial channels
 - Same for block devices and I/O

- Several constructors allow arbitrary layouts
 - Recursive specification possible
 - *Declarative* specification of data-layout
 - “what” and not “how”, leaves optimization to implementation (*many unexplored* possibilities!)
 - Choosing the right constructors is not always simple

更灵活 / 更复杂



Derived Datatype Example







MPI_Type_contiguous

```
MPI_Type_contiguous(int count, MPI_Datatype oldtype,  
MPI_Datatype *newtype)
```

- Contiguous array of *oldtype*
- *newtype* is the datatype obtained by concatenating *count* copies of *oldtype*







MPI_Type_vector

```
MPI_Type_vector(int count, int blocklength, int stride,  
MPI_Datatype oldtype, MPI_Datatype *newtype)
```

- Specify strided blocks of data of oldtype
- Very useful for Cartesian arrays



vector



struct

struct

struct

struct


vector



MPI_Type_create_hvector

```
MPI_Type_create_hvector(int count, int blocklength,  
                        MPI_Aint stride, MPI_Datatype oldtype,  
                        MPI_Datatype *newtype)
```

- Create **non-unit** strided vectors
- Useful for composition, e.g., vector of structs






MPI_Type_create_indexed_block

```
MPI_Type_create_indexed_block(int count, int blocklength,  
    int *array_of_displacements, MPI_Datatype  
    oldtype, MPI_Datatype *newtype)
```

- Pulling irregular subsets of data from a single array
 - dynamic codes with index lists, expensive though!
 - blen=2
 - displs={0,5,8,13,18}





MPI_Type_indexed

```
MPI_Type_indexed(int count, int *array_of_blocklengths,  
                 int *array_of_displacements, MPI_Datatype  
                 oldtype, MPI_Datatype *newtype)
```

- Like indexed_block, but can have different block lengths
 - `blen={1,1,2,1,2,1}`
 - `displs={0,3,5,9,13,17}`






MPI_Type_create_struct

```
MPI_Type_create_struct(int count, int array_of_blocklengths[],  
                      MPI_Aint array_of_displacements[],  
                      MPI_Datatype array_of_types[],  
                      MPI_Datatype *newtype)
```

- Most general constructor, allows different types and arbitrary arrays (also most costly)





MPI_Type_create_subarray

```
MPI_Type_create_subarray(int ndims, int array_of_sizes[],  
                        int array_of_subsizes[], int array_of_starts[],  
                        int order, MPI_Datatype oldtype,  
                        MPI_Datatype *newtype)
```

- Convenience function for creating datatypes for array segments
- Specify subarray of n-dimensional array (sizes) by start (starts) and size (subsize)

(0,0)	(0,1)	(0,2)	(0,3)
(1,0)	(1,1)	(1,2)	(1,3)
(2,0)	(2,1)	(2,2)	(2,3)
(3,0)	(3,1)	(3,2)	(3,3)



MPI_BOTTOM and MPI_Get_address

■ MPI_BOTTOM

- ☐ is the absolute zero address
- ☐ Portability (e.g., may be non-zero in globally shared memory)

■ MPI_Get_address

- ☐ Returns address relative to MPI_BOTTOM
- ☐ Portability (do not use “&” operator in C!)

■ Very important to

- ☐ build struct datatypes
- ☐ If data spans multiple arrays



Commit, Free, and Dup

■ **MPI_Type_commit**

- ☐ Types must be committed before use
- ☐ Only the ones that are used!
- ☐ may perform heavy optimizations (and will hopefully)

■ **MPI_Type_free**

- ☐ Free MPI resources of datatypes
- ☐ Does not affect types built from it

■ **MPI_Type_dup**

- ☐ Duplicates a type
- ☐ Library abstraction (composability)



Datatype Selection Order

- Simple and effective performance model:
 - More parameters == slower

predefined < contig < vector < index_block < index < struct

- Some (most) MPIs are inconsistent
 - But this rule is portable
- Advice to users:
 - Try datatype “compression” bottom-up

W. Gropp et al.: Performance Expectations and Guidelines for MPI Derived Datatypes



Code Example - 4

- *code/test4/*
 - `mpi_dedrived_datatype_contiguous.c`
 - `mpi_dedrived_datatype_index.c`
 - `mpi_dedrived_datatype_struct.c`
 - `mpi_dedrived_datatype_vector.c`



Content

- Non-blocking communication
- Group (collective) communication
- MPI Datatypes
- [Virtual Topology](#)



MPI Virtual Topologies

- In terms of MPI, a **virtual topology** describes a mapping/ordering of MPI processes into a geometric "shape".
- Virtual topologies are built upon MPI communicators and groups.
 - It is an attribute of processes only in the group.



MPI Virtual Topologies

- MPI topologies are **virtual**:
 - The term “Virtual Topology” gives this main idea: **machine independent**
 - no relation between the physical structure of the parallel machine and the process topology.
- Must be "programmed" by the application developer.



Why use Virtual Topologies

■ Convenience:

- Virtual topologies may be useful for applications with **specific communication patterns** - patterns that match an MPI topology structure.
- For example, a Cartesian topology might prove convenient for an application that requires 4-way nearest neighbor communications for grid based data.



Why use Virtual Topologies

■ Communication Efficiency :

- Some hardware architectures may impose penalties for communications between successively distant "nodes".
- A particular implementation may optimize process mapping based upon the physical characteristics of a given parallel machine.
- The mapping of processes into an MPI virtual topology is dependent upon the MPI implementation, and may be totally ignored.




MPI Topology History


- Convenience functions (in MPI-1)
 - Useful especially for Cartesian topologies
 - Query neighbors in n-dimensional space
 - Graph topology: each rank specifies full graph ☹
- Scalable Graph topology (MPI-2.2)
 - Graph topology: each rank specifies its neighbors **or** an arbitrary subset of the graph
- Neighborhood collectives (MPI-3.0)
 - Adding communication functions defined on graph topologies (neighborhood of distance one)

Types of Virtual Topologies

- There are two types of MPI topologies:



Cartesian Topology



Graph Topology



1 - Cartesian virtual topology

```
MPI_Cart_create(MPI_Comm comm_old, int ndims, const int *dims,  
                const int *periods, int reorder, MPI_Comm *comm_cart)
```

- Specify **ndims**-dimensional topology
 - ☐ Optionally **periodic** in each dimension (Torus)
- Some processes may return MPI_COMM_NULL
 - ☐ Product sum of dims must be $\leq P$
- **Reorder** argument allows for topology mapping
 - ☐ Each calling process may have a new rank in the created communicator
 - ☐ Data has to be remapped manually



MPI_Cart_create Example

```
int dims[3] = {5,5,5};  
int periods[3] = {1,1,1};  
MPI_Comm topocomm;  
MPI_Cart_create(comm, 3, dims, periods, 0, &topocomm);
```

- Creates logical 3-d Torus of size 5x5x5
- But we're starting MPI processes with a one-dimensional argument (-p X)
 - User has to determine size of each dimension
 - Often as “square” as possible, MPI can help!



MPI_Dims_create

```
MPI_Dims_create(int nnodes, int ndims, int *dims)
```

- Create dims array for Cart_create with *nnodes* and *ndims*
 - Dimensions are as close as possible (well, in theory)
- Non-zero entries in dims will not be changed
 - nnodes must be multiple of all non-zeroes



MPI_Dims_create Example

```
int p;  
MPI_Comm_size(MPI_COMM_WORLD, &p);  
MPI_Dims_create(p, 3, dims);  
  
int periods[3] = {1,1,1};  
MPI_Comm topocomm;  
MPI_Cart_create(comm, 3, dims, periods, 0, &topocomm);
```

- Makes life a little bit easier
 - Some problems may be better with a non-square layout though



Cartesian Query Functions

- Library support and convenience!
 - `MPI_Cartdim_get()`
 - Gets dimensions of a Cartesian communicator
 - `MPI_Cart_get()`
 - Gets size of dimensions
 - `MPI_Cart_rank()`
 - Translate coordinates to rank
 - `MPI_Cart_coords()`
 - Translate rank to coordinates



Cartesian Communication Helpers

```
MPI_Cart_shift(MPI_Comm comm, int direction, int disp,  
               int *rank_source, int *rank_dest)
```

- Shift in one dimension
 - Dimensions are numbered from 0 to ndims-1
 - Displacement indicates neighbor distance (-1, 1, ...)
 - May return MPI_PROC_NULL
- Very convenient, all you need for nearest neighbor communication
 - No “over the edge” though




Cartesian Example

- Code/test5/
 - mpi_topology_cartesian.c
- A simplified mapping of processes into a Cartesian virtual topology (Grid)

0 (0,0)	1 (0,1)	2 (0,2)	3 (0,3)
4 (1,0)	5 (1,1)	6 (1,2)	7 (1,3)
8 (2,0)	9 (2,1)	10 (2,2)	11 (2,3)
12 (3,0)	13 (3,1)	14 (3,2)	15 (3,3)

2 - Graph Topology

- More generally, the process organizing is described by a graph.
- Elements of Graph Topology:
 - Communication link
 - Nodes in the graph
 - Neighbours of per node
 - Type of mapping






Graph Topology

■ Elements of Graph Topology:

- Nodes:
 - Processors
- Lines:
 - Communicators between nodes
- Arrows:
 - Show origins and destinations of links
- Index:
 - array of integers describing node degrees

Node	Neighbors	Edges
0	1	1
1	2	0, 2
2	2	1, 3
3	1	2





Distributed Graph

- *MPI_Graph_create* is discouraged
 - Not scalable
 - Not deprecated yet but hopefully soon
- **New distributed interface:**
 - Scalable, allows distributed graph specification
 - Either local neighbors **or** any edge in the graph
 - Specify edge weights
 - Meaning undefined but optimization opportunity for vendors!
 - Info arguments
 - Communicate assertions of semantics to the MPI library
 - E.g., semantics of edge weights




MPI_Dist_graph_create_adjacent

```
MPI_Dist_graph_create_adjacent(MPI_Comm comm_old,  
                               int indegree, const int sources[], const int sourceweights[],  
                               int outdegree, const int destinations[], const int destweights[],  
                               MPI_Info info, int reorder, MPI_Comm *comm_dist_graph)
```

- indegree, sources, ~weights – source proc. Spec.
- outdegree, destinations, ~weights – dest. proc. spec.
- info, reorder, comm_dist_graph – as usual
- directed graph
- Each edge is specified twice, once as out-edge (at the source) and once as in-edge (at the dest)

MPI_Dist_graph_create_adjacent

- Process 0:
 - Indegree: 0
 - Outdegree: 2
 - Dests: {3,1}
- Process 1:
 - Indegree: 3
 - Outdegree: 2
 - Sources: {4,0,2}
 - Dests: {3,4}
- ...





MPI_Dist_graph_create

```
MPI_Dist_graph_create(MPI_Comm comm_old, int n,  
    const int sources[], const int degrees[],  
    const int destinations[], const int weights[],  
    MPI_Info info, int reorder, MPI_Comm *comm_dist_graph)
```

- **n** – number of source nodes
- **sources** – n source nodes
- **degrees** – number of edges for each source
- **destinations, weights** – dest. processor specification
- **info, reorder** – as usual
- **More flexible and convenient**
 - Requires global communication
 - Slightly more expensive than adjacent specification



MPI_Dist_graph_create


■ Process 0:

- N: 2
- Sources: {0,1}
- Degrees: {2,1}^{*}
- Dests: {3,1,4}

■ Process 1:

- N: 2
- Sources: {2,3}
- Degrees: {1,1}
- Dests: {1,2}

■ ...



* Note that in this example, process 0 specifies only one of the two outgoing edges of process 1; the second outgoing edge needs to be specified by another process



Distributed Graph Neighbor Queries

```
MPI_Dist_graph_neighbors_count(  
    MPI_Comm comm,  
    int *indegree, int *outdegree, int *weighted)
```

- Query the number of neighbors of **calling process**
- Returns indegree and outdegree;
- Also info if weighted



Distributed Graph Neighbor Queries

```
MPI_Dist_graph_neighbors(MPI_Comm comm,  
                          int maxindegree, int sources[], int sourceweights[],  
                          int maxoutdegree, int destinations[], int destweights[])
```

- Query the neighbor list of **calling process**
- Optionally return weights



Further Graph Queries

`MPI_Topo_test(MPI_Comm comm, int *status)`

- Status is either:
 - `MPI_GRAPH` (ugs)
 - `MPI_CART`
 - `MPI_DIST_GRAPH`
 - `MPI_UNDEFINED` (no topology)
- Enables to write libraries on top of MPI topologies!



Neighborhood Collectives

- Topologies implement **no communication!**
 - Just helper functions
- Collective communications only cover some patterns
 - E.g., no stencil pattern
- Several requests for “build your own collective” functionality in MPI
 - Neighborhood collectives are a simplified version
 - Cf. Datatypes for communication patterns!





Cartesian Neighborhood Collectives

- Communicate with direct neighbors in Cartesian topology
 - Corresponds to `cart_shift` with `disp=1`
 - Collective (all processes in comm must call it, including processes without neighbors)
 - Buffers are laid out as neighbor sequence:
 - Defined by order of dimensions, first negative, then positive
 - $2 \times \text{ndims}$ sources and destinations
 - Processes at borders (`MPI_PROC_NULL`) leave holes in buffers (will not be updated or communicated)!



Cartesian Neighborhood Collectives

- Buffer ordering example:





Graph Neighborhood Collectives

- Collective Communication along arbitrary neighborhoods
 - Order is determined by order of neighbors as returned by `(dist_)graph_neighbors`.
 - Distributed graph is directed, may have different numbers of send/recv neighbors
 - Can express dense collective operations 😊
 - Any persistent communication pattern!



MPI_Neighbor_allgather

```
MPI_Neighbor_allgather(  
    const void* sendbuf, int sendcount, MPI_Datatype sendtype,  
    void* recvbuf, int recvcount, MPI_Datatype recvtpe,  
    MPI_Comm comm)
```

- Sends the same message to all neighbors
- Receives indegree distinct messages
- Similar to MPI_Gather
 - The all prefix expresses that each process is a “root” of his neighborhood
- Vector version for full flexibility



MPI_Neighbor_alltoall

```
MPI_Neighbor_alltoall(  
    const void* sendbuf, int sendcount, MPI_Datatype sendtype,  
    void* recvbuf, int recvcount, MPI_Datatype recvttype,  
    MPI_Comm comm)
```

- Sends outdegree distinct messages
- Received indegree distinct messages
- Similar to MPI_Alltoall
 - Neighborhood specifies full communication relationship
- Vector and w versions for full flexibility




Nonblocking Neighborhood Collectives


```
MPI_Ineighbor_allgather(..., MPI_Request *req);  
MPI_Ineighbor_alltoall(..., MPI_Request *req);
```

- Very similar to **nonblocking** collectives
- Collective invocation
- Matching in-order (no tags)
 - No wild tricks with neighborhoods! In order matching per communicator!


Examples of Graph Topology



(a) Random



(b) Ring



(c) Wheel



Example topologies

MPI_Dist_graph_create_adjacent:

Topology/Process	indegree	sources	outdegree	destinations
random/0	3	1,3,4	2	2,3
random/1	2	2,4	2	0,4
random/2	1	0	2	1,3
random/3	2	0,2	2	0,4
random/4	2	1,3	2	0,1
ring/ i ($0 \leq i \leq 4$)	1	$i - 1 \bmod 5$	1	$i + 1 \bmod 5$
wheel/0	0	-	4	1,2,3,4
wheel/1	2	0,4	1	2
wheel/2	2	0,1	1	3
wheel/3	2	0,2	1	4
wheel/4	2	0,3	1	1



Example topologies

MPI_Dist_graph_create:

Topology/Process	n	sources	degrees	destinations
random/0	1	2	1	1
random/1	1	1	2	0,4
random/2	2	0,4	2,1	2,3,1
random/3	2	4,3	1,1	0,0
random/4	2	2,3	1,1	3,4
ring/ i ($0 \leq i \leq 4$)	1	$i - 1 \bmod 5$	1	i
wheel/0	1	2	1	3
wheel/1	1	3	1	4
wheel/2	1	4	1	1
wheel/3	1	1	1	2
wheel/4	1	0	4	1,2,3,4



Graph Example

- Code/test5/
 - `mpi_graph_topology_adjacent.cpp`
 - `mpi_graph_topology.cpp`
- Reference:
 - The Scalable Process Topology Interface of MPI 2.2



references

- Parallel Programming with MPI, Argonne National Laboratory, <http://www.anl.gov/events/parallel-programming-mpi>.
- A Comprehensive MPI Tutorial Resource, <https://github.com/wesleykendall/mpitutorial>.
- 张武生等, MPI并行程序设计实例教程, 清华大学出版社。
- 迟学斌等, 并行计算与实现技术, 科学出版社。