


# Introduction to CUDA

## *(5) Performance Considerations*

# Reference

- [CUDA C Programming Guide](https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html),
  - <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>
- **Programming Massively Parallel Processors,**
  - A Hands-on Approach
  - Third Edition
  - Chapter 5



# Resources and Performance

- Performance of a parallel program can vary greatly depending on the **resource constraints** of the computing hardware.
- In different applications, **different constraints** may dominate and become *bottlenecks*.
- It is important to find:
  - the **major types of resource constraints** in a CUDA device;
  - how they can affect the kernel execution performance.

# Content

- Global Memory Bandwidth
- More on Memory Parallelism
- Warps and SIMD Hardware
- Dynamic partitioning of resources

# Global Memory Bandwidth

- Ideal




- Reality



# DRAM core arrays are slow

- Reading from a cell in the core array is a very slow process




- 10s of nanoseconds v.s. sub-nanosecond clock cycle
- Like determine the flavor of cup of coffee by smelling far away

# DRAM Burst



Each time a DRAM location is accessed, a range of consecutive locations that includes the requested location are actually accessed.

These consecutive locations accessed and delivered are referred to as DRAM bursts.



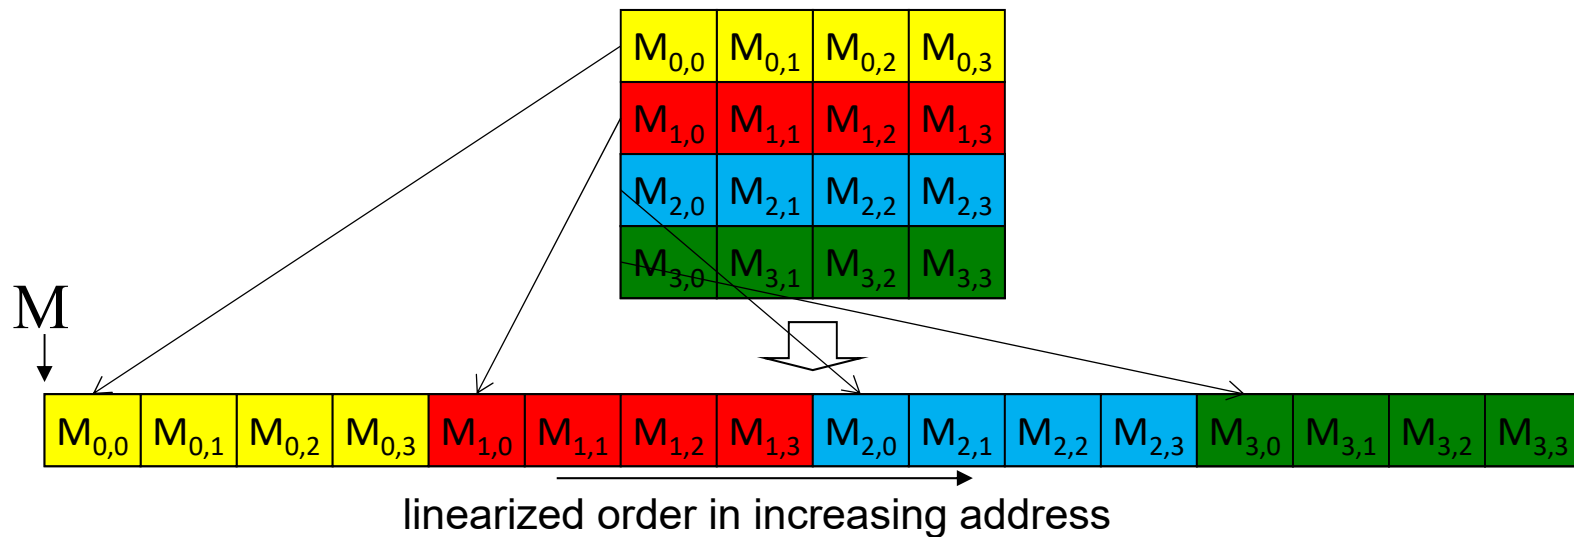
# Coalesce

- Threads in a warp execute the same instruction at any given time.
- Modern DRAM systems are designed to be always accessed in burst mode.
- The favorable access pattern is achieved When all threads in a warp access consecutive global memory locations.
- In this case, the hardware combines, or **coalesces**, all these accesses into a consolidated access to consecutive DRAM locations (burst).




# A 2D C Array in Linear Memory Space

- How the memory addresses are formed in accessing C multidimensional array elements?




# DRAM Burst – A System View




- Each address space is partitioned into burst sections
  - Whenever a location is accessed, all other locations in the same section are also delivered to the processor
- Basic example: a 16-byte address space, 4-byte burst sections
  - In practice, we have at least 4GB address space, burst section sizes of 128-bytes or more

# Memory Coalescing



- When all threads of a warp execute a load instruction, **if all accessed locations fall into the same burst section**, only one DRAM request will be made and the **access is fully coalesced**.

# Un-coalesced Accesses



- When the accessed locations spread across burst section boundaries:
  - Coalescing fails
  - Multiple DRAM requests are made

# A Simple Matrix Multiplication Kernel (review)

```
__global__
void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
    // Calculate the row index of the P element and M
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
    // Calculate the column index of P and N
    int Col = blockIdx.x * blockDim.x + threadIdx.x;

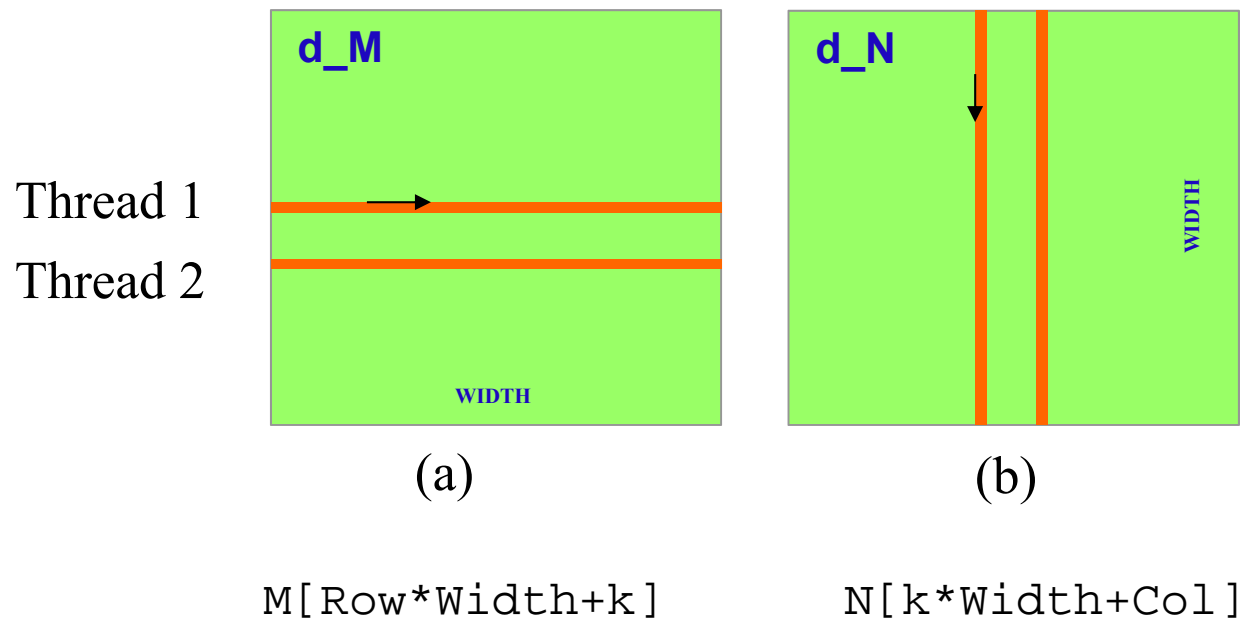
    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += M[Row*Width+k] * N[k*Width+Col];

        P[Row*Width+Col] = Pvalue;
    }
}
```

# Two Access Patterns

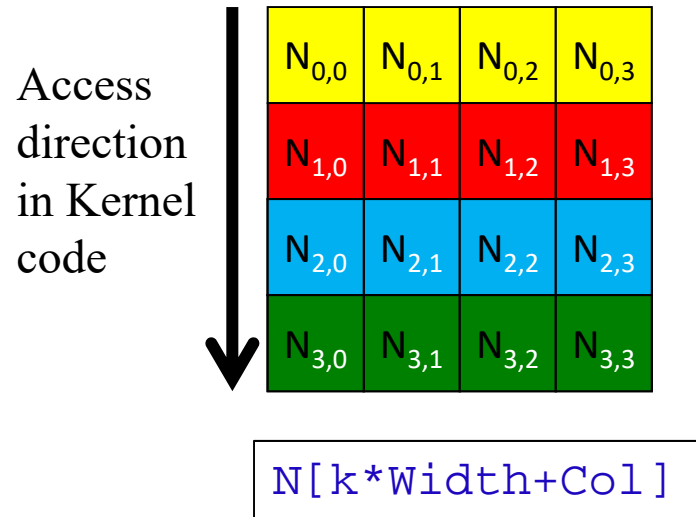
For M, threads in a warp read adjacent rows;

For N, threads in a warp read adjacent columns;



$k$  is loop counter in the inner product loop of the kernel code

# N accesses are coalesced.



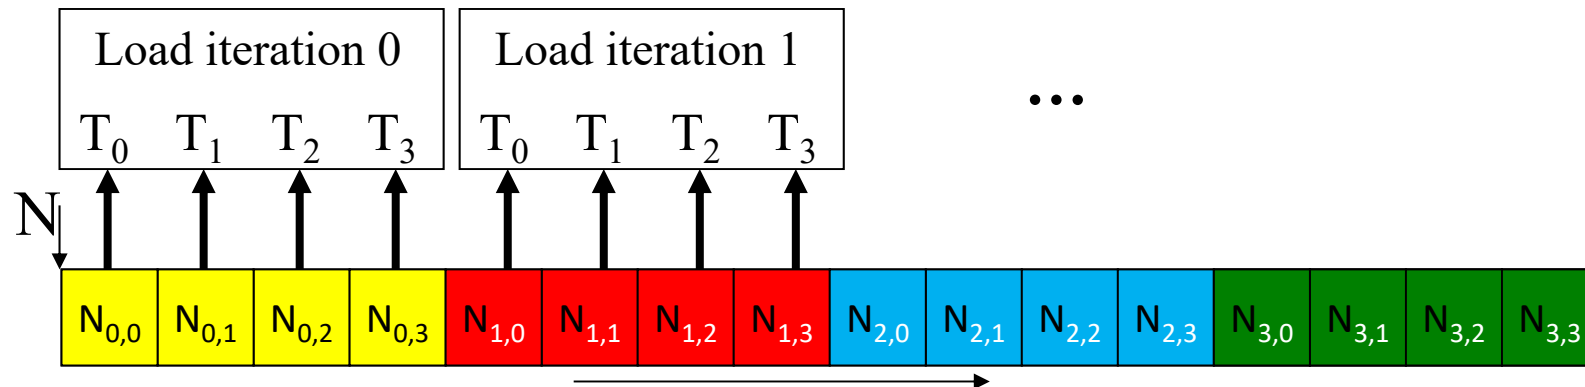
$k*Width$ :

same across all threads within a given iteration of the k loop.

Col:


$Col = blockIdx.x * blockDim.x + threadIdx.x$

$blockIdx.x$  and  $blockDim.x$  are same for all threads in the same block.



# M accesses are not coalesced.

Access  
direction  
in Kernel  
code



$M[\text{Row} * \text{Width} + k]$

k & Width:


same across all threads within a given iteration of the k loop.

Row:

$\text{Row} = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$

$\text{blockIdx.x}$  and  $\text{blockDim.x}$  are same for all threads in the same block.



...







# Use shared memory to enable coalescing in tiled matrix multiplication

Original  
Access  
Pattern




Corner Turning

Tiled  
Access  
Pattern



Copy into  
scratchpad  
memory



Perform  
multiplication  
with scratchpad  
values

# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width){  
1.  __shared__ float Mds[TILE_WIDTH][TILE_WIDTH];  
2.  __shared__ float Nds[TILE_WIDTH][TILE_WIDTH];  
3.  int bx = blockIdx.x;  int by = blockIdx.y;  
4.  int tx = threadIdx.x; int ty = threadIdx.y;  
    // Identify the row and column of the P element to work on  
5.  int Row = by * TILE_WIDTH + ty;  
6.  int Col = bx * TILE_WIDTH + tx;  
7.  float Pvalue = 0;  
    // Loop over the M and N tiles required to compute the P element  
8.  for (int ph = 0; ph < Width/TILE_WIDTH; ++ph) {  
    // Collaborative loading of M and N tiles into shared memory  
9.    Mds[ph][ph] = M[                ?                ];  
10.   Nds[ph][ph] = N[                ?                ];  
11.   __syncthreads();  
12.   for (int k = 0; k < TILE_WIDTH; ++k)  
13.       Pvalue += Mds[ty][k] * Nds[k][tx];  
14.   __syncthreads();  
15. }  
16. P[Row*Width+Col] = Pvalue;  
}
```

# Load M elements

- The linearized index calculation is equivalent to the two-dimensional array access expression:
- $M[\text{Row}][\text{ph} * \text{TILE\_WIDTH} + \text{tx}]$ .
  - column index used by the threads only differs in terms of **threadIdx**.
  - row index is determined by **blockIdx.y** and **threadIdx.y**
  - threads in the same thread block have **identical** **blockIdx.y/threadIdx.y** and **adjacent** **threadIdx.x**
- $M[\text{ty}][\text{tx}] = M[\text{Row} * \text{Width} + \text{ph} * \text{TILE\_WIDTH} + \text{tx}]$

# Load N elements

- The linearized index calculation is equivalent to the two-dimensional array access expression:
- $N[ph * TILE\_WIDTH + ty][bx * TILE\_WIDTH + tx]$ .
  - the row index  $ph * TILE\_WIDTH + ty$  has the same value for all threads with the same  $threadIdx.y$  value.
  - the column index calculation for each thread,  $Col = bx * TILE\_WIDTH + tx$ , where  $bx * TILE\_WIDTH$  is the same for all threads in the same block
- $Nds[ty][tx] = N[(ph * TILE\_WIDTH + ty) * Width + Col];$

# Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width){  
1.  __shared__ float Mds[TILE_WIDTH][TILE_WIDTH];  
2.  __shared__ float Nds[TILE_WIDTH][TILE_WIDTH];  
3.  int bx = blockIdx.x;  int by = blockIdx.y;  
4.  int tx = threadIdx.x; int ty = threadIdx.y;  
    // Identify the row and column of the P element to work on  
5.  int Row = by * TILE_WIDTH + ty;  
6.  int Col = bx * TILE_WIDTH + tx;  
7.  float Pvalue = 0;  
    // Loop over the M and N tiles required to compute the P element  
8.  for (int ph = 0; ph < Width/TILE_WIDTH; ++ph) {  
    // Collaborative loading of M and N tiles into shared memory  
9.      Mds[ty][tx] = M[Row*Width + ph*TILE_WIDTH + tx];  
10.     Nds[ty][tx] = N[(ph*TILE_WIDTH + ty)*Width + Col];  
11.     __syncthreads();  
12.     for (int k = 0; k < TILE_WIDTH; ++k)  
13.         Pvalue += Mds[ty][k] * Nds[k][tx];  
14.     __syncthreads();  
15. }  
16. P[Row*Width+Col] = Pvalue;  
}
```


~30x

# Content

- Global Memory Bandwidth
- More on Memory Parallelism
- Warps and SIMD Hardware
- Dynamic partitioning of resources

# Bursting is not sufficient

- DRAM systems typically employ two more forms of parallel organization:
  - banks and channels.
  - At the highest level, a processor contains one or more channels.
  - Each channel is a memory controller with a bus that connects a set of DRAM banks to the processor.




# Number of Channels

- The data transfer bandwidth of a bus is defined by its width and clock frequency:
  - For example, a 64-bit DDR bus with a clock frequency of 1 GHz has a bandwidth of  $8B * 2 * 1 \text{ GHz} = 16 \text{ GB/sec}$
- This seems to be a large number but is often too small for modern CPUs and GPUs.
  - CPU  $\rightarrow 32\text{GB/s} \rightarrow 2 \text{ channels}$
  - GPU  $\rightarrow 128\text{GB/s} \rightarrow 8 \text{ channels}$



# Number of Banks

- For each channel, the number of banks is determined by fully utilizing the bandwidth of the bus:



(A) Single-Bank burst (dark segment) timing, dead time (light segment) on interface




Bank 0 bursts

Bank 1 bursts

(B) Two-Bank burst timing, reduced dead time

# Multiple DRAM Banks

In general, if the ratio of the cell array access latency and data transfer time is  $R$ , we need to have at least  $R+1$  banks.



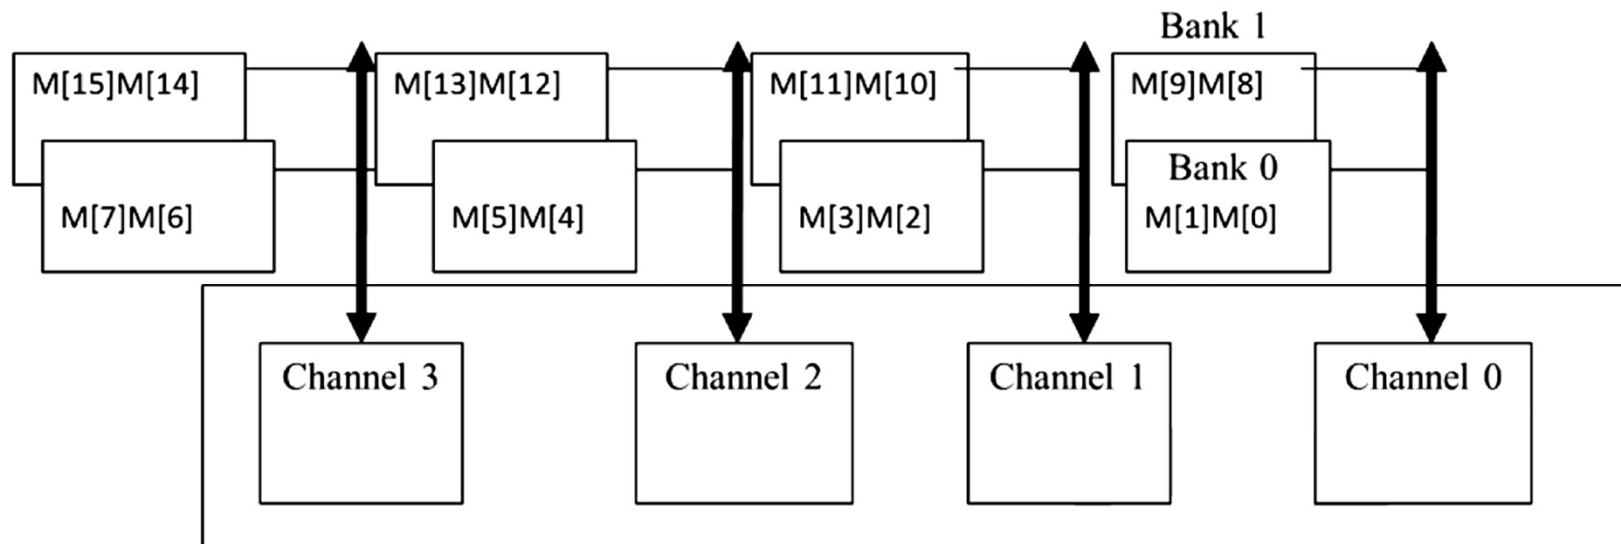
- Reasons to have more banks:
  - Avoid bank conflicts;
  - Manufacturability;

# Data Distribution

- There is an important connection between the parallel execution of threads and the parallel organization of the DRAM system:
  1. sufficient number of threads making simultaneous memory accesses.
  2. these memory accesses must be evenly **distributed to the channels and banks**.
  3. each access to a bank must also be a **coalesced access**.

# Data Distribution

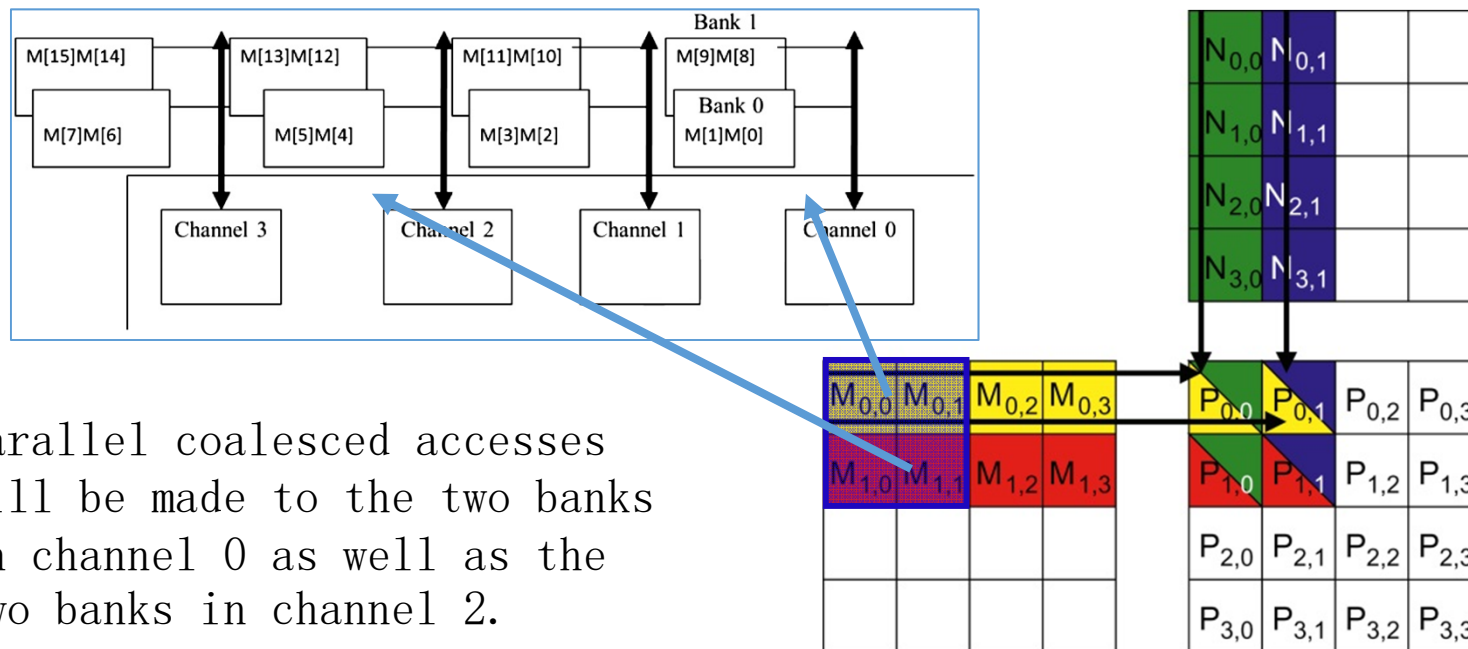
- The distribution scheme *spreads* the elements across the banks and channels in the system:



*Interleaved data distribution*


# Interleaved Data Distribution

- Example:
  - multiplication will be performed with  $2 \times 2$  thread blocks and  $2 \times 2$  tiles



# Interleaved Data Distribution

- Example:




Tiles loaded by	Block 0,0	Block 0,1	Block 1,0	Block 1,1
Phase 0 (2D index)	M[0][0], M[0][1], M[1][0], M[1][1]	M[0][0], M[0][1], M[1][0], M[1][1]	M[2][0], M[2][1], M[3][0], M[3][1]	M[2][0], M[2][1], M[3][0], M[3][1]
Phase 0 (linearized index)	M[0], M[1], M[4], M[5]	M[0], M[1], M[4], M[5]	M[8], M[9], M[12], M[13]	M[8], M[9], M[12], M[13]
	cache			
Phase 1 (2D index)	M[0][2], M[0][3], M[1][2], M[1][3]	M[0][2], M[0][3], M[1][2], M[1][3]	M[2][2], M[2][3], M[3][2], M[3][3]	M[2][2], M[2][3], M[3][2], M[3][3]
Phase 1 (linearized index)	M[2], M[3], M[6], M[7]	M[2], M[3], M[6], M[7]	M[10], M[11], M[14], M[15]	M[10], M[11], M[14], M[15]

Multiplying  $8 \times 8$  matrices with the same  $2 \times 2$  thread block configuration:  
What's the number of channels to make use?

# Content

- Global Memory Bandwidth
- More on Memory Parallelism
- Warps and SIMD Hardware
- Dynamic partitioning of resources

# Warps as Scheduling Units




- Each block is divided into 32-thread warps
  - An implementation technique
  - Warps are scheduling units in SM
  - Threads in a warp execute in Single Instruction Multiple Data (SIMD) manner
  - The number of threads in a warp may vary in future generations



# Warps in Multi-dimensional Thread Blocks

- The thread blocks are first linearized into 1D in row major order **before partitioning into warps**.
  - In x-dimension first, y-dimension next, and z-dimension last
  - For this example, all 16 threads form half a warp.
  - The warp will be padded with another 16 threads to complete a 32-thread warp.



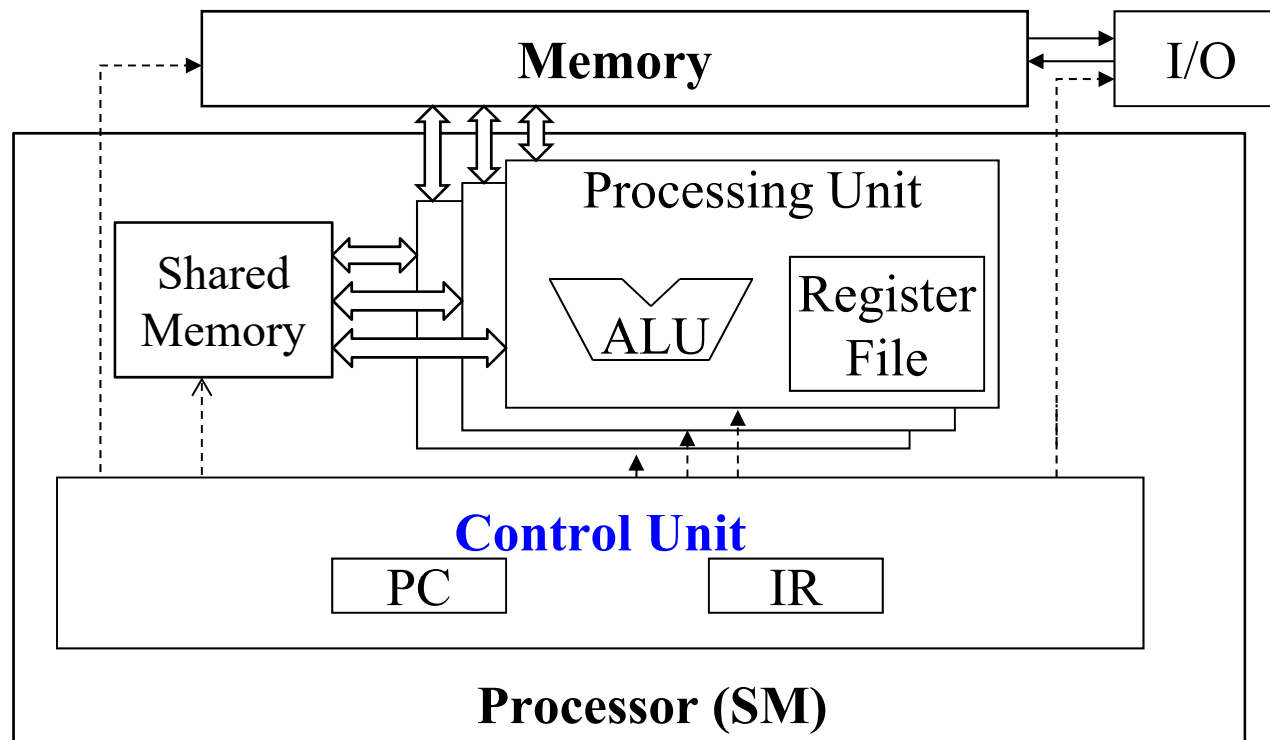
$T_{y,x}$ ,  $x$  being `threadIdx.x` and  $y$  being `threadIdx.y`.

# Blocks are partitioned after linearization

- Linearized thread blocks are partitioned
  - Thread indices within a warp are consecutive and increasing
  - Warp 0 starts with Thread 0
- Partitioning scheme is consistent across devices
  - Thus you can use this knowledge in control flow
  - However, the exact size of warps may change from generation to generation
- DO NOT rely on any ordering within or between warps
  - If there are any dependencies between threads, you must `__syncthreads()` to get correct results (more later).

# SMs are SIMD Processors

- Control unit for instruction fetch, decode, and control is **shared among multiple processing units**
  - Control overhead is minimized



# SIMD Execution Among Threads in a Warp

- All threads in a warp must execute the same instruction at any point in time
- This works efficiently if all threads follow the same control flow path
  - All if-then-else statements make the same decision
  - All loops iterate the same number of times

# Control Divergence

- Control divergence occurs when threads in a warp take different control flow paths by making different control decisions
  - Some take the then-path and others take the else-path of an if-statement
  - Some threads take different number of loop iterations than others
- The execution of threads taking different paths are serialized in current GPUs
  - The control paths taken by the threads in a warp are traversed one at a time until there is no more.
  - During the execution of each path, all threads taking that path will be executed in parallel
  - The number of different paths can be large when considering nested control flow statements

# Control Divergence Examples

- Divergence can arise when branch or loop condition is a function of thread indices:
  - Example kernel statement with divergence:
    - `if (threadIdx.x > 2) { }`
    - This creates two different control paths for threads in a block
    - Decision granularity < warp size; threads 0, 1 and 2 follow different path than the rest of the threads in the first warp
  - Example without divergence:
    - `if (blockIdx.x > 2) { }`
    - Decision granularity is a multiple of blocks size;
    - all threads in any given warp follow the same path.

# Example: Vector Addition Kernel

## Device Code

```
// Compute vector sum  $C = A + B$   
// Each thread performs one pair-wise addition  
  
__global__  
void vecAddKernel(float* A, float* B, float* C,  
    int n)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    if(i < n) C[i] = A[i] + B[i];  
}
```

# Analysis for vector size of 1,000 elements

- Assume that block size is 256 threads
  - 8 warps in each block
- All threads in Blocks 0, 1, and 2 are within valid range
  - i values from 0 to 767
  - There are 24 warps in these three blocks, none will have control divergence
- Most warps in Block 3 will not control divergence
  - Threads in the warps 0-6 are all within valid range, thus no control divergence
- One warp in Block 3 will have control divergence
  - Threads with i values 992-999 will all be within valid range
  - Threads with i values of 1000-1023 will be outside valid range
- Effect of serialization on control divergence will be small
  - 1 out of 32 warps has control divergence
  - The impact on performance will likely be less than 3%



# Performance Impact of Control Divergence


- Boundary condition checks are vital for complete functionality and robustness of parallel code
  - The **tiled matrix multiplication** kernel has many boundary condition checks
  - The concern is that these checks may cause significant performance degradation
  - For example, see the tile loading code below:

```
if(Row < Width && t * TILE_WIDTH+tx < Width) {  
    ds_M[ty][tx] = M[Row * Width + p * TILE_WIDTH + tx];  
} else {  
    ds_M[ty][tx] = 0.0;  
}
```

```
if (p*TILE_WIDTH+ty < Width && Col < Width) {  
    ds_N[ty][tx] = N[(p*TILE_WIDTH + ty) * Width + Col];  
} else {  
    ds_N[ty][tx] = 0.0;  
}
```

# Two types of blocks in loading M Tiles

- 1. Blocks whose tiles are all within valid range until the last phase.
- 2. Blocks whose tiles are partially outside the valid range all the way




# Analysis of Control Divergence Impact

- Assume 16x16 tiles and thread blocks
- Each thread block has 8 warps (256/32)
- Assume square matrices of 100x100
- Each thread will go through 7 phases (ceiling of  $100/16$ )
- There are 49 thread blocks (7 in each dimension)

# Control Divergence in Loading M Tiles


- Assume 16x16 tiles and thread blocks
- Each thread block has 8 warps (256/32)
- Assume square matrices of 100x100
- Each warp will go through 7 phases (ceiling of  $100/16$ )
- There are 42 ( $6*7$ ) Type 1 blocks, with a total of 336 ( $8*42$ ) warps
- They all have 7 phases, so there are 2,352 ( $336*7$ ) warp-phases
- The warps have control divergence only in their last phase
- 336 warp-phases have control divergence



# Control Divergence in Loading M Tiles (Type 2)


Type 2:

- the 7 block assigned to load the bottom tiles, with a total of 56 ( $8*7$ ) warps
- They all have 7 phases, so there are 392 ( $56*7$ ) warp-phases
- The first 2 warps in each Type 2 block will stay within the valid range until the last phase
- The 6 remaining warps stay outside the valid range
- So, only 14 ( $2*7$ ) warp-phases have control divergence



# Overall Impact of Control Divergence

- Type 1 Blocks: 336 out of 2,352 warp-phases have control divergence
- Type 2 Blocks: 14 out of 392 warp-phases have control divergence
- The performance impact is expected to be less than 12% ( $350/2,944$  or  $(336+14)/(2352+14)$ )



# Additional Comments

- The estimated performance impact is data dependent.
  - For larger matrices, the impact will be significantly smaller
- In general, the impact of control divergence for boundary condition checking for large input data sets should be insignificant
  - One should not hesitate to use boundary checks to ensure full functionality
- The fact that a kernel is full of control flow constructs does not mean that there will be heavy occurrence of control divergence

# Content

- Global Memory Bandwidth
- More on Memory Parallelism
- Warps and SIMD Hardware
- Dynamic partitioning of resources



# Execution resources

- The execution resources in an SM include:
  - registers,
  - shared memory,
  - thread block slots,
  - and thread slots.
- These resources are **dynamically partitioned** and assigned to threads to support their execution.

# Execution resources

- **Thread slots:**

- These thread slots are partitioned and assigned to thread blocks during runtime.

- **For example: Fermi devices have 1536 thread slots.**

- If each thread block consists of 512 threads, the 1536 thread slots are partitioned and assigned to three blocks.
  - In this case, each SM can accommodate up to three thread blocks due to limitations on thread slots.
  - If each thread block contains 256 threads, the 1536 thread slots are partitioned and assigned to 6 thread blocks.

- **The ability to dynamically partition the thread slots among thread blocks makes SMs versatile.**

# Execution resources

- **Block slots:**
  - subtle interactions between block slots and thread slots.
- For example: Fermi devices have 1536 thread slots.
  - If each block has 128 threads, the 1536 thread slots can be partitioned and assigned to 12 blocks.
  - However, since there are only 8 block slots in each SM, only 8 blocks will be allowed.
  - This means that in the end, only 1024 of the thread slots will be utilized.
- Therefore, to fully utilize both the block slots and thread slots, one needs at least 256 threads in each block.

# Execution resources

- **Registers:**

- the automatic variables declared in a CUDA kernel are placed into registers.

- For example: assume each SM has 16,384 registers.

- If the kernel uses 10 registers per thread for  $16 \times 16$  thread blocks, how many threads can run on each SM?  $10 * 16 * 16 = 2560$
- The number of registers required by six blocks is 15,360. OK
- Adding another block would require 17,920 registers? NOT
- How about additional two automatic variables?
- $12 * 16 * 16 = 3072$ ; registers required by six blocks is now 18,432. NOT

- The CUDA runtime will reduce the number of blocks assigned to each SM by one.

- This, however, reduces the number of threads running on an SM from 1536 to 1280.
- “performance cliff”

# Execution resources

- **Shared memory:**
  - another resource that is dynamically partitioned at run-time.
- For example: Tiled algorithms
  - require a large amount of shared memory to be effective.
  - large shared memory usage can reduce the number of thread blocks running on an SM.
- Reduced thread parallelism can negatively affect the utilization of the memory access bandwidth;
  - reduced memory access throughput, in turn, can further reduce the thread execution throughput.

# Execution resources

- Constraints of all the dynamically partitioned resources interact with each other in a complex manner.
- **CUDA Occupancy Calculator**
  - a downloadable Excel sheet
  - calculates the actual number of threads running on each SM for a particular device given the usage of resources by a kernel.

# summary

- The memory bandwidth is a first-order performance factor in a massively parallel processor
  - DRAM bursts, banks, and channels
  - All concepts are also applicable to modern multicore processors
- The memory coalescing is important for effectively utilizing memory bandwidth in CUDA
  - Its origin in DRAM burst
  - Checking if a CUDA memory access is coalesced
  - Techniques for improving memory coalescing in CUDA code